


Network Names in Content-Centric Networking


Cesar Ghali, Gene Tsudik and Christopher Wood
University of California, Irvine
{cghali, gene.tsudik, woodc1}@uci.edu

CCN Names

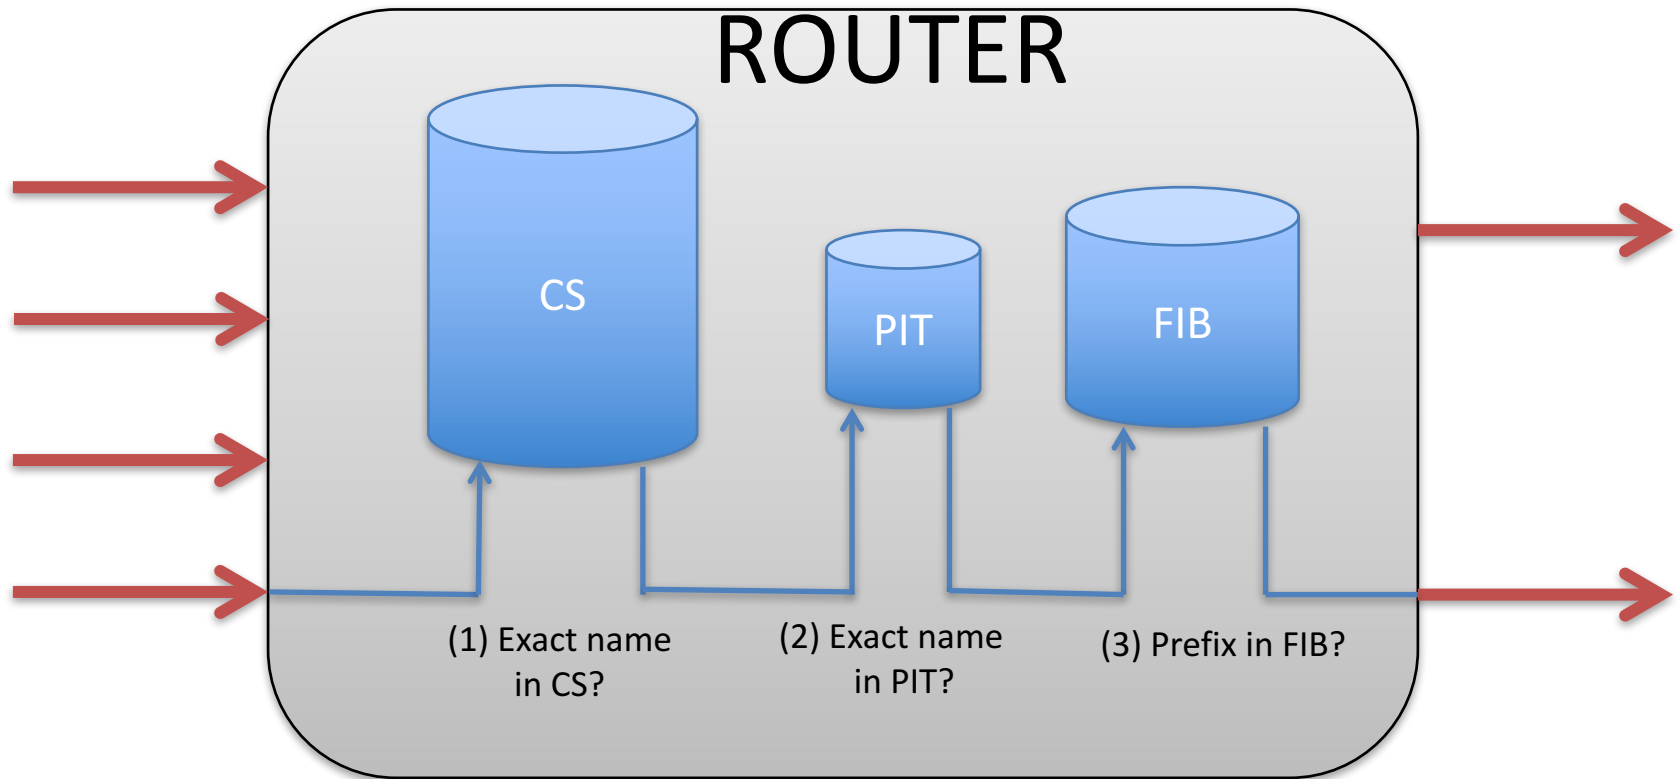
- Expressed as URIs
 - /a/b/foo
 - /us/edu/uci/cs/tsudik/papers/acm-icn16.pdf
- Encoded as TLVs



Names in CCN Applications



Names in the Network



Dual Roles for Names

- Applications use names to meaningfully express and identify content
 - Human readability is nice!
- Network entities (routers) use names as sequences of binary strings
 - A router doesn't care about readability or arbitrarily long components

Q: Why is the same representation used at both layers & in both contexts?

Outline

- CCN Network Names
- Name Translation and Its Implications
- Translation Analysis
- Related & Future Work

CCN Network Names

Cryptographic Hash Function?

Goal: translate application names into a format that:

- Facilitates standard network processing (exact match and LPM searches)
- Removes arbitrarily long names from the network
- Removes all location-irrelevant information from the name (as seen by routers)

Translation  should:

be a **bijection** or very close to one

map **arbitrarily long strings** to **fixed-length output**

only apply to **location-specific parts of a name.**

Name Translation Example

/a/b/foo/version=0x00/chunk=0x01/PID=0x02
payload: <bar>



Application-specific
components not
Contained in FIBs

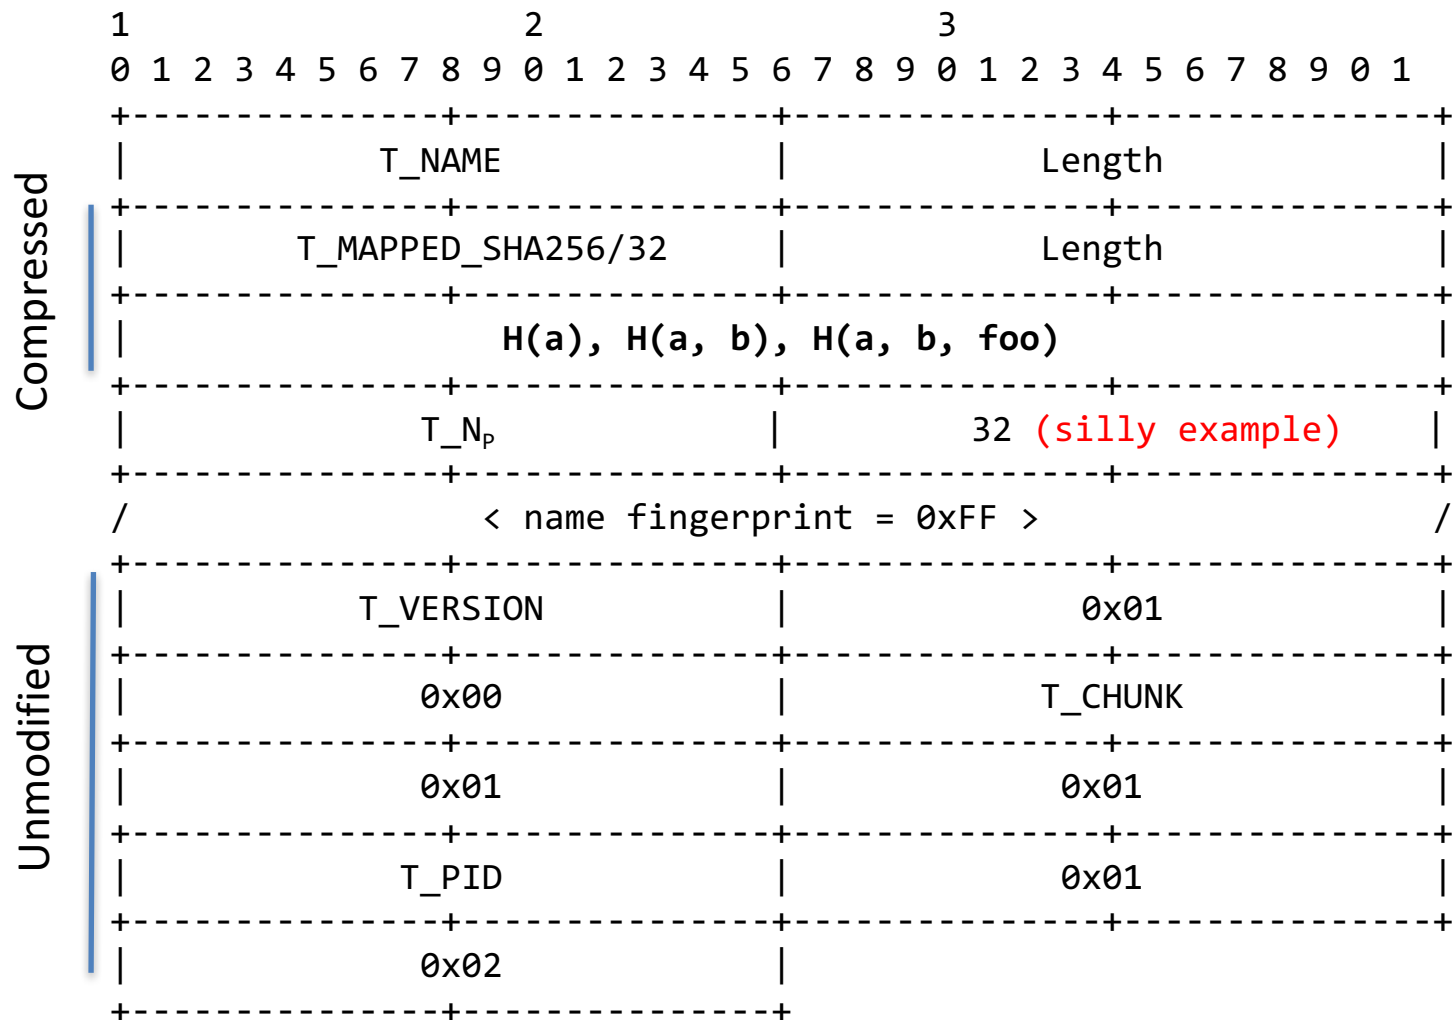
Unique Name
Fingerprint

T_NAME locator components
possibly contained in FIBs

/H(a)/H(a,b)/H(a,b,foo)/version=0x00/chunk=0x01/PID=0x02/**N_p=0xFF**
payload: <bar>

$$N_p = H'(\textbf{/a/b/foo/version=0x00/chunk=0x01/PID=0x02})$$

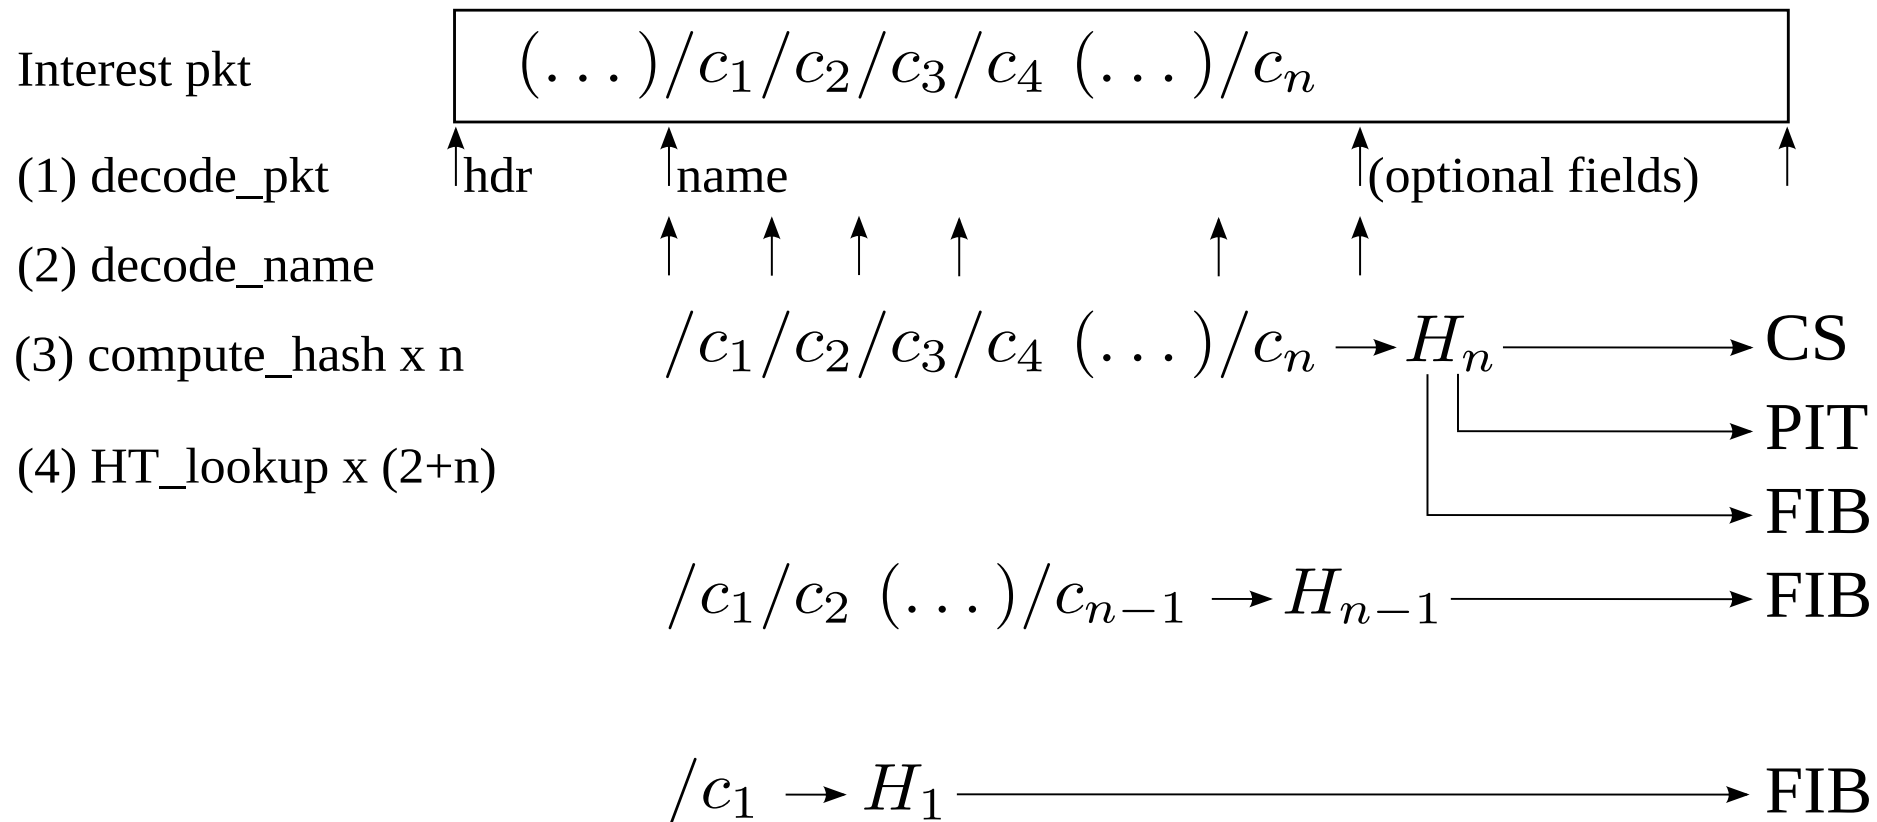
Name Encoding



Application Processing

- Consumer:
 - Maps application names to network names for outgoing interests
 - Inverts mapping for incoming content
- Producer:
 - Pre-computes network names for all its locator prefixes
 - Stores in “inversion table”
 - Looks up content corresponding to incoming interests based on this inversion table
 - Signed content objects contain \mathbf{N}_p
 - might also carry app name

Current Network Processing



**We obviate the need to hash in order to perform:
FIB, CS and PIT lookups**

Processing Summary

| Entity: | Impact: |
|----------|---|
| Consumer | Increased online processing |
| Producer | Increased offline computation & storage of inversion table |
| Router | Faster FIB lookup with pre-computed name-prefix hashes Faster PIT and CS lookups with N_p (Potential) benefits due to fixed-size N_p |

Analysis Setup

Questions:

- How big should a hash digest be?
 - What is the impact on packet sizes?
 - What are the resulting collision properties?
- What's consumer processing overhead?

Unibas URI data set: <http://www.icn-names.net/>



Name Properties

| Names | | Name segments | | Segments per Name | |
|--------------------------------|-------------|-----------------------------------|---------------|--------------------------------------|---------------|
| Number of names | 870'896'633 | Total number of segments | 4'855'203'042 | Total number of segments | 4'855'203'042 |
| Average name length (bytes) | 57.95 | Average segment length (bytes) | 10.39 | Average segments per name | 5.57 |
| Name length standard deviation | 77.60 | Segment length standard deviation | 30.02 | Segments per name standard deviation | 8.14 |
| Minimum name length (bytes) | 1 | Minimum segments length (bytes) | 1 | Minimum segments per name | 1 |
| Maximum name length (bytes) | 764'867 | Maximum segments length (bytes) | 764'867 | Maximum segments per name | 210'658 |

Name Properties


| Number of segments n | Number of names | Percentage |
|------------------------|--------------------|----------------|
| 1 | 13'952 | 0.002% |
| 2 | 141'904 | 0.016% |
| 3 | 71'327'647 | 8.190% |
| 4 | 187'307'048 | 21.507% |
| 5 | 253'852'565 | 29.148% |
| 6 | 144'130'578 | 16.550% |
| 7 | 93'837'904 | 10.775% |
| 8 | 70'875'144 | 8.138% |
| 9 | 25'611'959 | 2.941% |
| 10 | 10'464'092 | 1.202% |
| 11 | 3'973'961 | 0.456% |
| 12 | 4'546'842 | 0.522% |
| 13 | 1'206'905 | 0.139% |
| 14 | 835'124 | 0.096% |
| 15 | 844'552 | 0.097% |
| 16 | 195'491 | 0.022% |
| 17 | 121'486 | 0.014% |
| 18 | 317'628 | 0.036% |
| 19 | 168'228 | 0.019% |
| 20 | 50'742 | 0.006% |
| Total | 869'823'752 | 99.876% |

Name Size Impact (for interests)





Ratio of “network name” to “standard name”
as size of T()/H() grows

N_p Size Overhead (for Content)



Collision Assessment (32-bit hash)



Processing Overhead (consumer)

Per-name costs (μ s):

- Average: **1,029.279** (≈ 1 ms)
- Minimum: 3.812
- Maximum: 2,474.567

→ *Reasonable compared to network I/O*

Throughput (c/b):

- Average: **1,577.688**
- Minimum: 1,218.037
- Maximum: 3,494.538

Experimental setting:

- Intel 2.8 GHz Core i7
- Un-optimized implementation based on PARC Libparc libraries

Ideal setting:

- Use Intel intrinsics for hashing (~ 9 c/b)
- Work on wire-encoded packets

Wrap-up

- Motivated separating CCN application and network names
- A concrete mechanism for name translation function
- Assessed quality of name translation function and performance implications for all CCN entities

Related Work*

- ❑ CCN names
 - ❑ Requirements [Ghodsi et al., ICN'11]
 - ❑ Location-agnostic names [Van Adrichem et al., Nomen'13]
- ❑ Focus on FIB algorithm improvements
 - ❑ FIB algorithm modifications based on tries, hash tables, Bloom Filters, etc.
[Quan et al., Networking'13], [Perino et al., ANCS'14],
[So et al., ANCS'13], [Fukushima et al., Nomen'13]
 - ❑ Several rely on lexicographical name ordering
 - ❑ Our scheme breaks this!
 - ❑ Hop-by-hop optimizations, e.g., passing length of previously matched name in FIB

... but not on **FIB inputs**

**See paper for references*

Future Work

- Compare performance of current FIB techniques with and without network names
- Play with various hash functions & sizes
- Explore further uses for translation function $T()$

/this/is/the/end/version=0x00/chunk=0x01/PID=0x02