

Cryptographic Algorithms and Security Protocols for ICN

Christopher A. Wood

UC Irvine, PARC

woodc1@uci.edu, cwood@parc.com

Background


- Security was not a goal of the IP protocol
 - Added after the fact via IPSec, SSL/TLS, etc.
- Existing IP security solutions are:
 - Rooted in old (long-standing) security solutions
 - Change and adapt very slowly
 - Only now starting to look ahead (PFS, PQ algorithms and protocols, etc.)

Looking Forward

- ICN architectures are:
 - Built on a **clean slate**
 - Can use **new and modern crypto**


ICN Network Stack

An ICN Network Stack

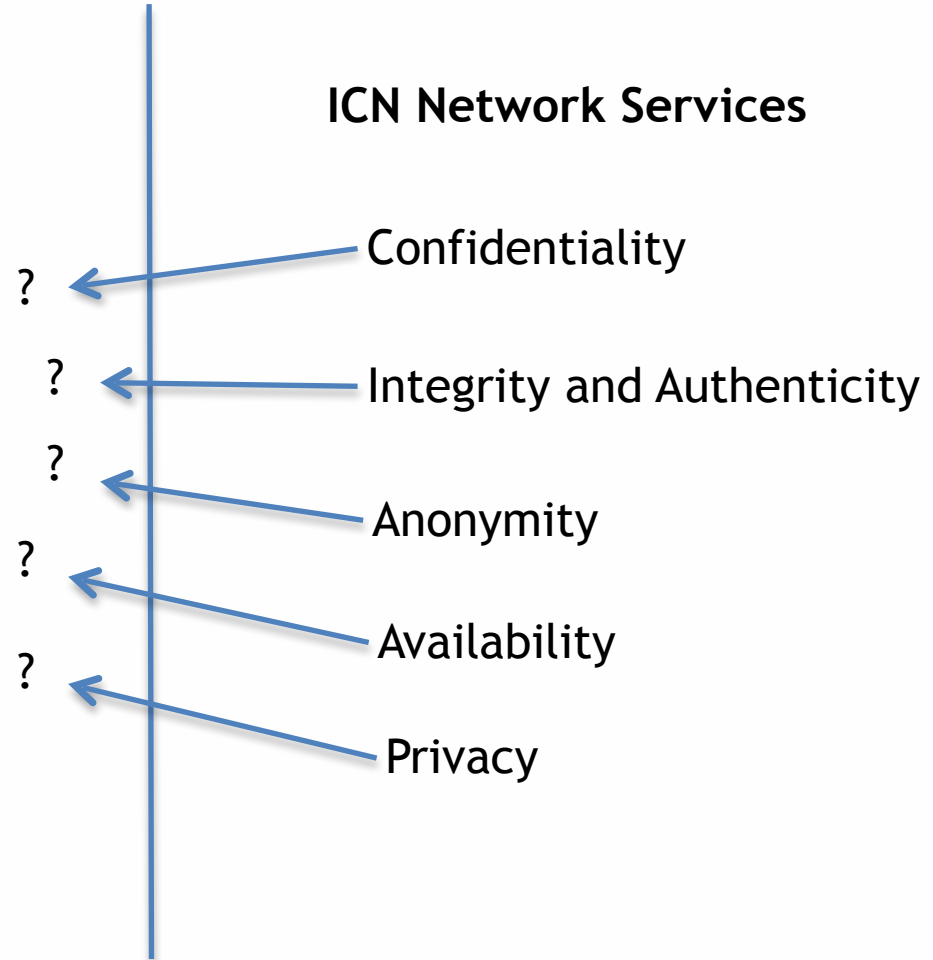


ICN Network Services

An ICN Network Stack




ICN Network Services




ICN Network Services

An ICN Network Stack



ICN Network Services



Q: What crypto and security techniques can enable these core services?

Selection Criteria

- Goal: Crypto that **can enables essential network services**
- What's essential?
 - Integrity and authenticity
 - Anonymity
 - Privacy (e.g., as per RFC 6973)
 - Availability
- What's non essential?
 - Confidentiality (application-layer concern)

Topic Breakdown


- Integrity and authenticity
 - Hash-based signatures
- Anonymity
 - TOR and Hornet
- Availability
 - Authenticated Denial of Existence (DoE)
- Privacy
 - Encrypted Deep Packet Inspection (DPI)
 - Password-Authenticated and Non-Interactive Key Exchange (PAKE and NIKE)
 - Private Information Retrieval (PIR)
 - Randomizable public-key encryption
 - Secure searchable encryption (SSE)
 - Oblivious data structures

Integrity and Authentication

Hash-Based Signatures

- Traditional signature schemes are based on trapdoor functions
 - RSA, DSA, ECDSA, etc.
 - PQ-secure?
- Hash-based signatures are quantum-secure, e.g., they don't fall to Shor's algorithm
- Based on one-time signatures (OTSs)
 - A key pair can be used only once

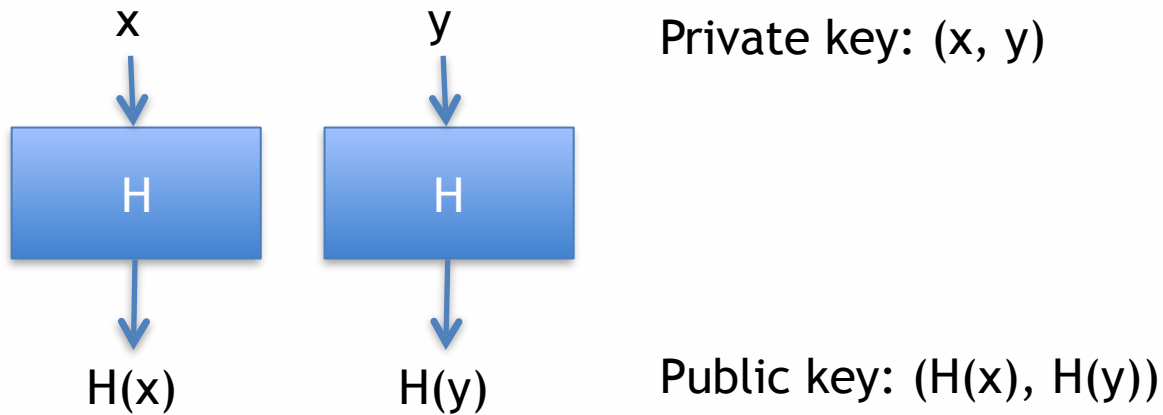
Lamport OTS Idea [x]



Private key: (x, y)


Public key: $(H(x), H(y))$

Lamport OTS Idea




To sign a '0': release x
To sign a '1': release y

Merkle Tree Idea




Merkle Tree Idea



To sign the first message m_1 : provide $(A, B, H(CD), \text{sign}_{k_A}(m_1))$

Merkle Tree Idea


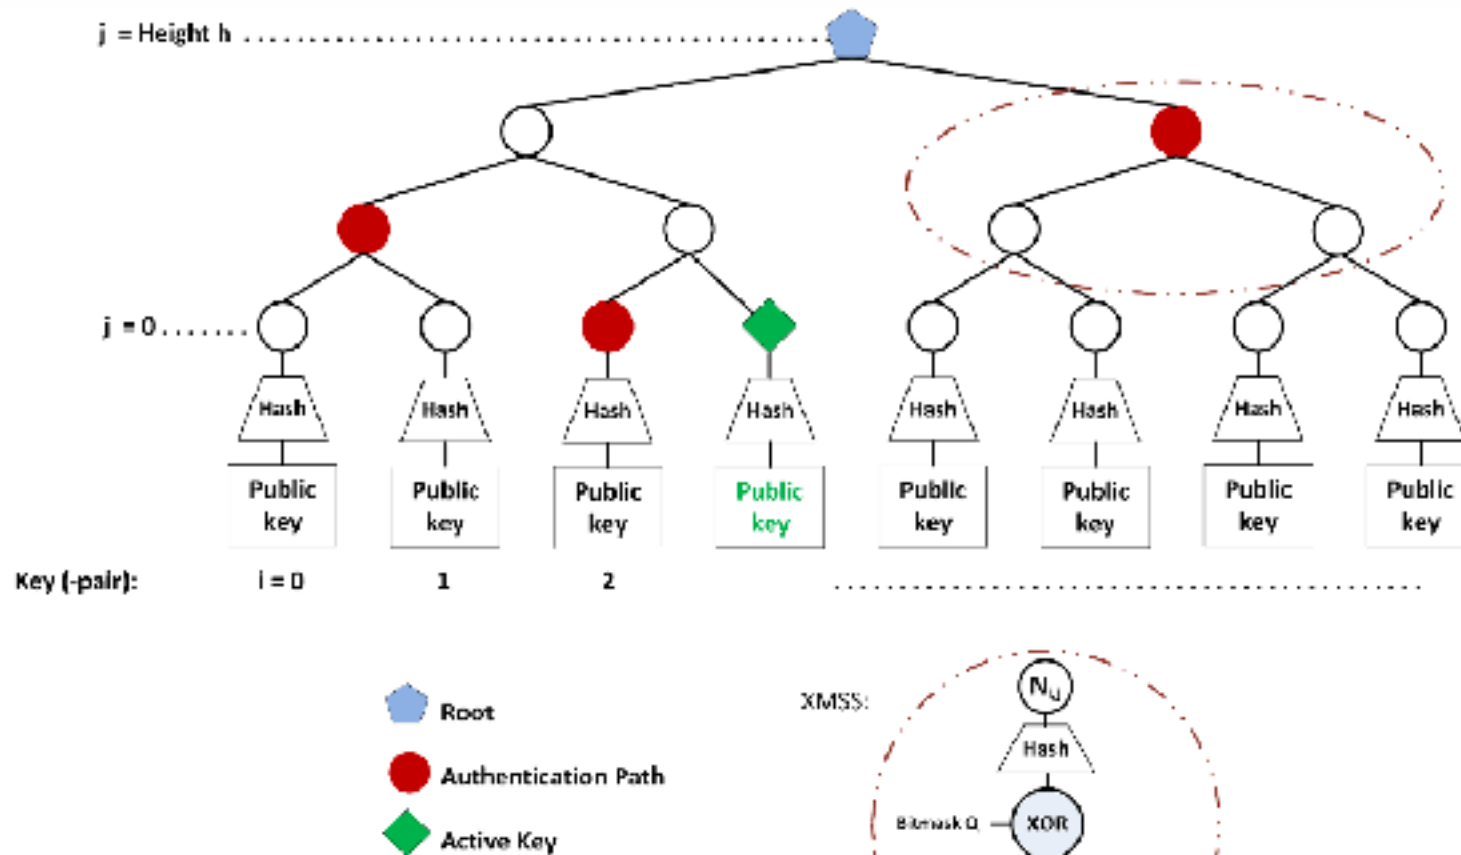


To sign the third message m_3 : provide $(H(AB), C, D, \text{sign}_{k_A}(m_3))$

Merkle Trees [x]

- The tree dimensions determine the number of messages that can be signed
- Other drawbacks:
 - It is secure in the Random Oracle model
 - The length is **proportional to the number of messages** that are signed
 - It is **stateful**
 - Store the index and intermediate tree results

XMSS Tree (Buchmann et al.)




Stateless Hash-Based Signature

- Similar to Merkle trees but the “index” is chosen at random
- Requires huge trees to avoid collisions
 - For security parameter $\lambda=128$ we require $2^{2\lambda}$ leaves (by birthday paradox)
- OTS secret keys are generated pseudorandomly

Overview

Signer chooses random $r \in \{2^{255}, 2^{255} + 1, \dots, 2^{256} - 1\}$,
 uses one-time public key T_r to sign message;
 uses one-time public key T_i to sign (T_{2i}, T_{2i+1}) for $i < 2^{255}$.
 Generates i th secret key as $H_k(i)$ where k is master secret.



SPINCS (Bernstein et al.)

- Stateless hash-based signature
 - Inspired by Goldreich stateless hash-based scheme [X]
- Dramatically reduces tree size
- Replace one-time leaves with “few-time leaves”
- Use XMSS-like per-node masks

Hash-Based Signature Recap

- When to use?
 - Long-term security in a post-quantum world
- Will ICN data packets live forever?

Anonymity

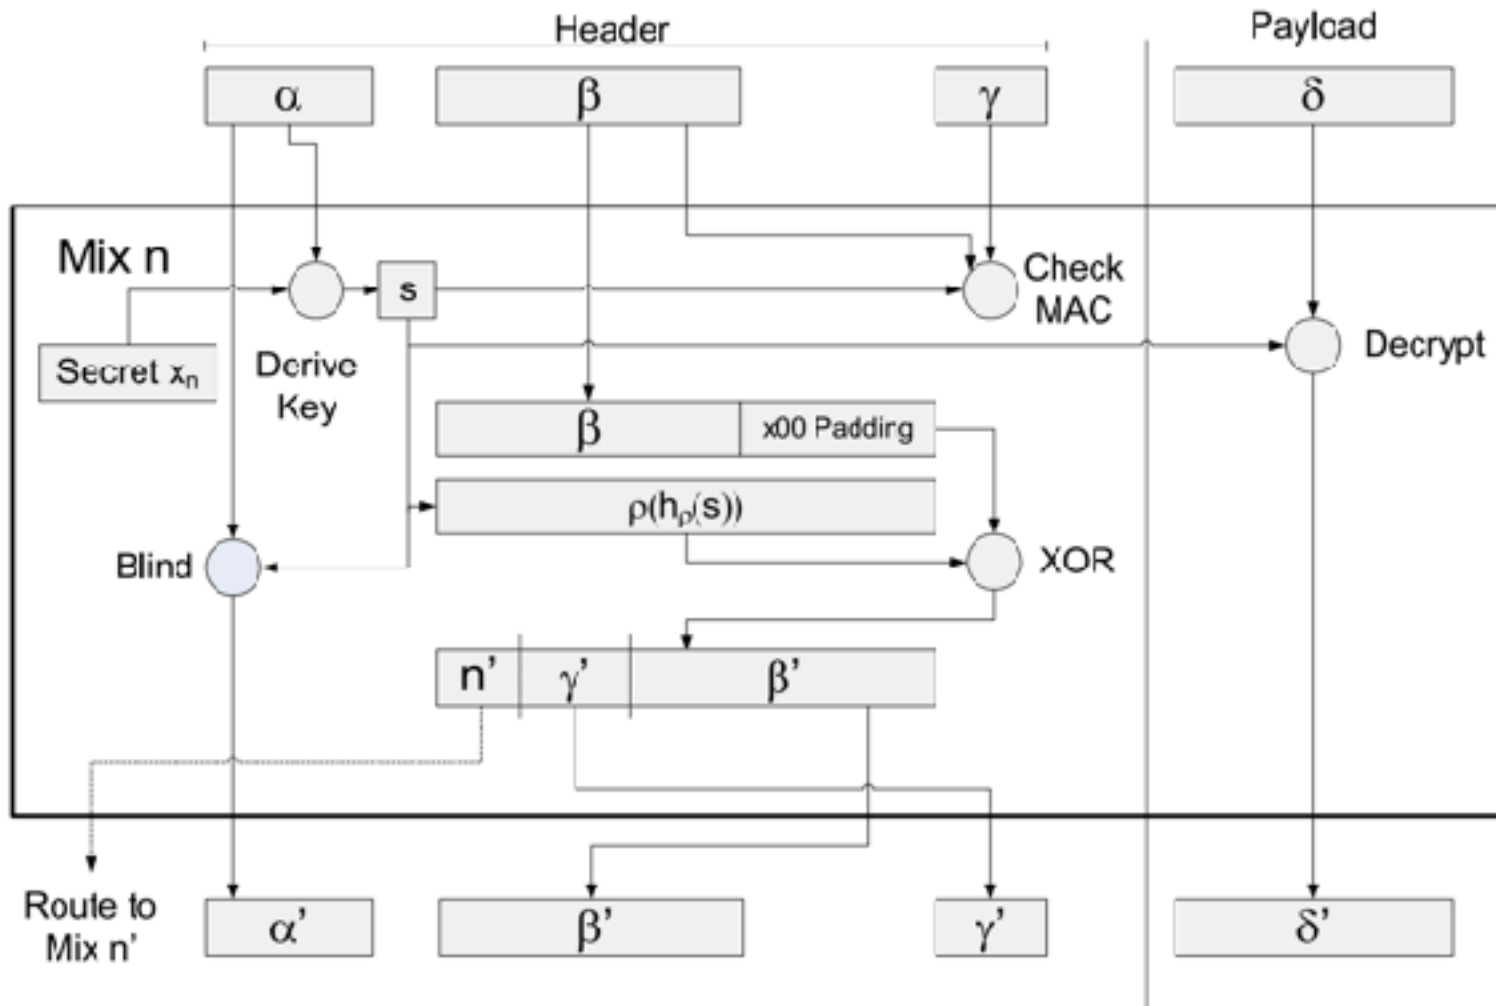
Onion Routing

- Based on Chaum's mixnets
- TOR (The Onion Router) is the most ubiquitous variant
 - Not provably secure
 - Performance could be improved

Sphinx Protocol [x]

- Provably secure mix protocol
- Mixture of public- and symmetric-key cryptography for each packet
 - Derive per-packet MAC and encryption keys for each anonymizer
 - Onion-encrypt the payload with a large-block-size block cipher
 - Compute the MAC carefully
 - Wrap the derivation seed using key encapsulation

Sphinx Protocol



HORNET Protocol [x]

- Built on Sphinx to establish symmetric keys with each anonymizer
 - Extended to include per-hop “forwarding segments” used to process packets in a circuit
- Data packets are created with onion-encrypted payloads and anonymous headers (ADHR)


HORNET Session Setup Packet

type	hops	EXP
Sphinx Header		
Sphinx Payload		
FS Payload		

HORNET Data Packet

type	hops	nonce
ADHR		
Data Payload		

HORNET Data Packet Internals



Onion Routing Recap

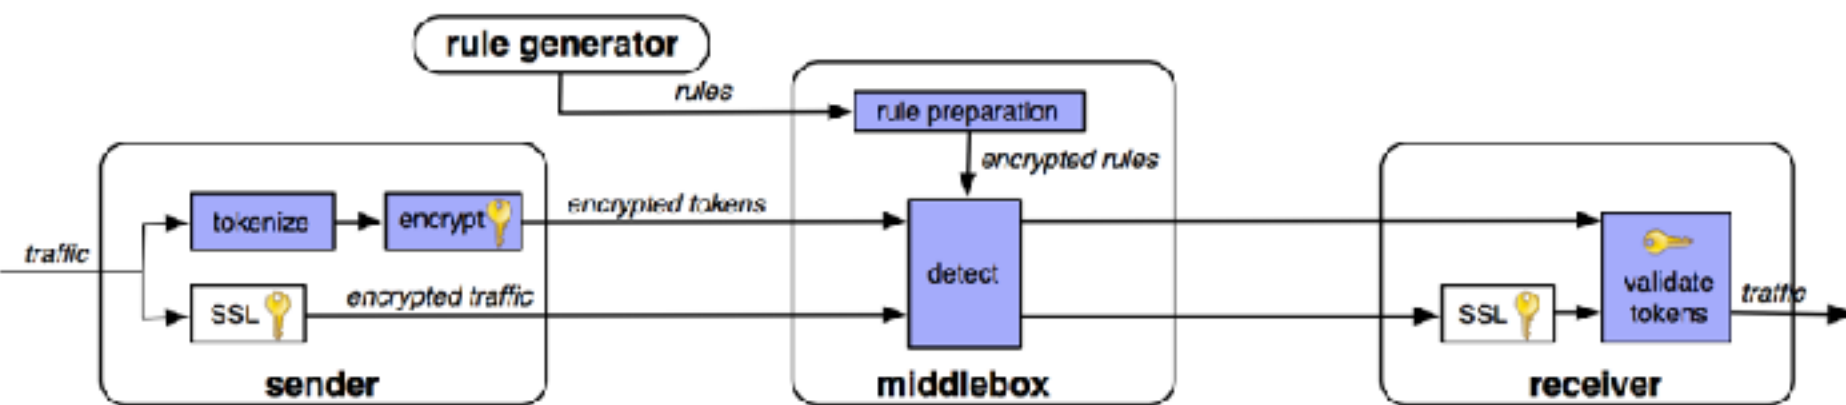
- ICN names reveal **too much**
- Use onion routing to onion-encrypt individual packets
 - G. Tsudik, E. Uzun, and C. A. Wood, AC3N: An API and Service for Anonymous Communication in Content-Centric Networking, in the Proceedings of CCNC 2016, Las Vegas, NV, USA, 2016.
 - S. DiBenedetto, P. Gasti, G. Tsudik and E. Uzun, ANDaNA: Anonymous Named Data Networking Application, the 19th Annual Network and Distributed System Security Symposium (NDSS), San Diego, CA, 2012.
- ... but do it right as a network service

Privacy

Encrypted DPI

- Perform deep-packet-inspection on encrypted packet payloads
 - Determine when a packet contains an encrypted version of a specific keyword
- Different measures of privacy
 - Exact-match privacy: only discover bytes that match target keywords
 - Probable cause privacy: decrypt a flow (entire packet) only if a keyword match is detected

BlindBox



BlindBox Details

- Packets are tokenized into tokens t_1, \dots, t_k
- For each token t , encrypt as follows
$$\text{salt}, \text{AES}_{\text{AES}-k(t)}(\text{salt}) \longrightarrow$$

ensures randomness of identical tokens
- To detect a token t , precompute salt and token combinations
- Speedup: use a single salt per token and then derive subsequent token encryptions based on a counter

Probable Cause Decryption

- Idea: embed message encryption key D_k in the encryption of each token

salt, $(\text{AES}_{\text{AES-}k(t)}(\text{salt}) \text{ XOR } D_k)$

BlindBox Rule Preparation

- The middlebox must learn $\text{AES-}k(t)$ without revealing the encryption key k
- Solution: Garbled circuit to compute
 $\text{AES-}k(t)$
for each token t

Privacy: Key Exchange

Password AKE

- Goal: create ephemeral keys from shared secrets – passwords (read: name) – in a way that is:
 - Not susceptible to offline dictionary attacks
 - Forward secure
- Most protocols are multi-round, e.g., J-PAKE
- Some protocols work in a **single round** without sacrificing forward secrecy

Non-Interactive Key Exchange (NIKE)

- Goal: Two parties with knowledge of each other's public keys agree on a shared secret without requiring any interaction [1]

Alice: x, g^x

Bob: y, g^y

Shared key: $H(\text{"Alice"}, \text{"Bob"}, g^{xy})$

- Use case: WSN shared key derivation
- Public-key encryption follows from NIKE

NIKE Protocols

- Setup: generate system parameters
- KeyGen: generate a private and public key pair for a given identity
- SharedKey: given (1) an identity, (2) its public key, (3) another identity, (4) and its secret key, compute and output a shared key

NIKE Protocols


- Setup: generate system parameters
- KeyGen: generate a private and public key pair for a given identity
- SharedKey: given (1) an **identity**, (2) its **public key**, (3) another identity, (4) and its secret key, compute and output a shared key



What's in an ICN certificate?...

Forward Secrecy?

- Yes – it's possible: evolve the public keys.
- Use multilinear maps and the tree-based key derivation technique
- Add an Update function to the scheme:
 - Move the secret key *forward* in time and space




Privacy: PIR

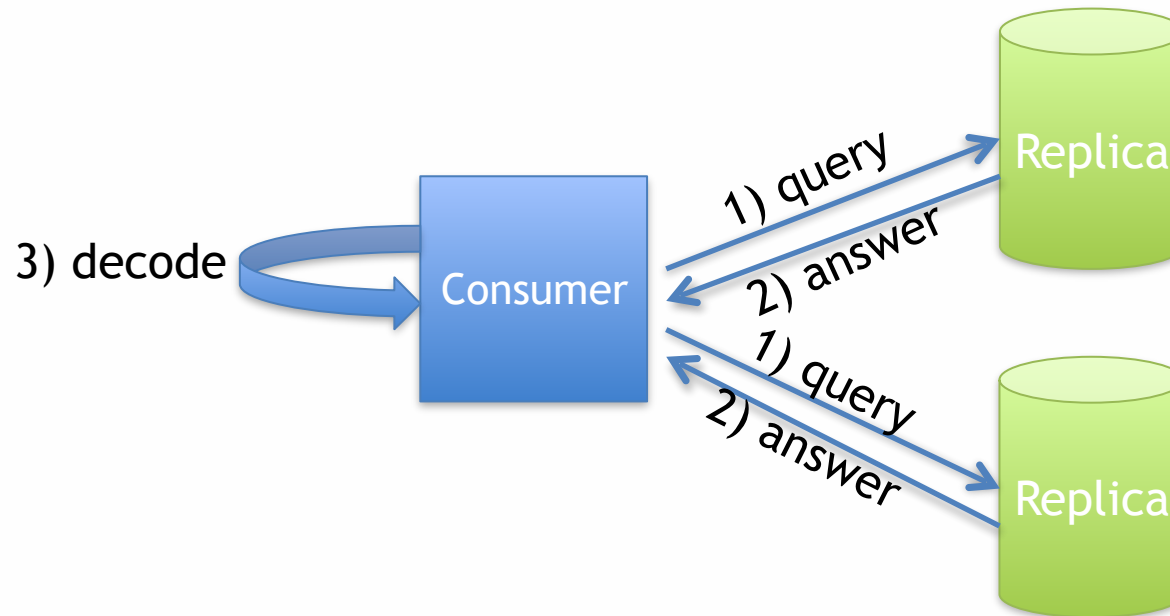
Private Information Retrieval

- A protocol to hide client requests to a server (e.g., a database)
- Two variants: computational (CPIR) and information-theoretic (IPIR)
- CPIR
 - Require only a single server
 - Much more computationally expensive
- IPIR
 - Require at least two non-colluding servers
 - More efficient but requires more communication

IPIR



IPIR



C. Tschudin, Private Information Retrieval over ICN,
INFOCOM 2016 NOM Workshop, April, 2016.


Privacy: Randomizable Encryption

Forward-Secure Public Key Encryption

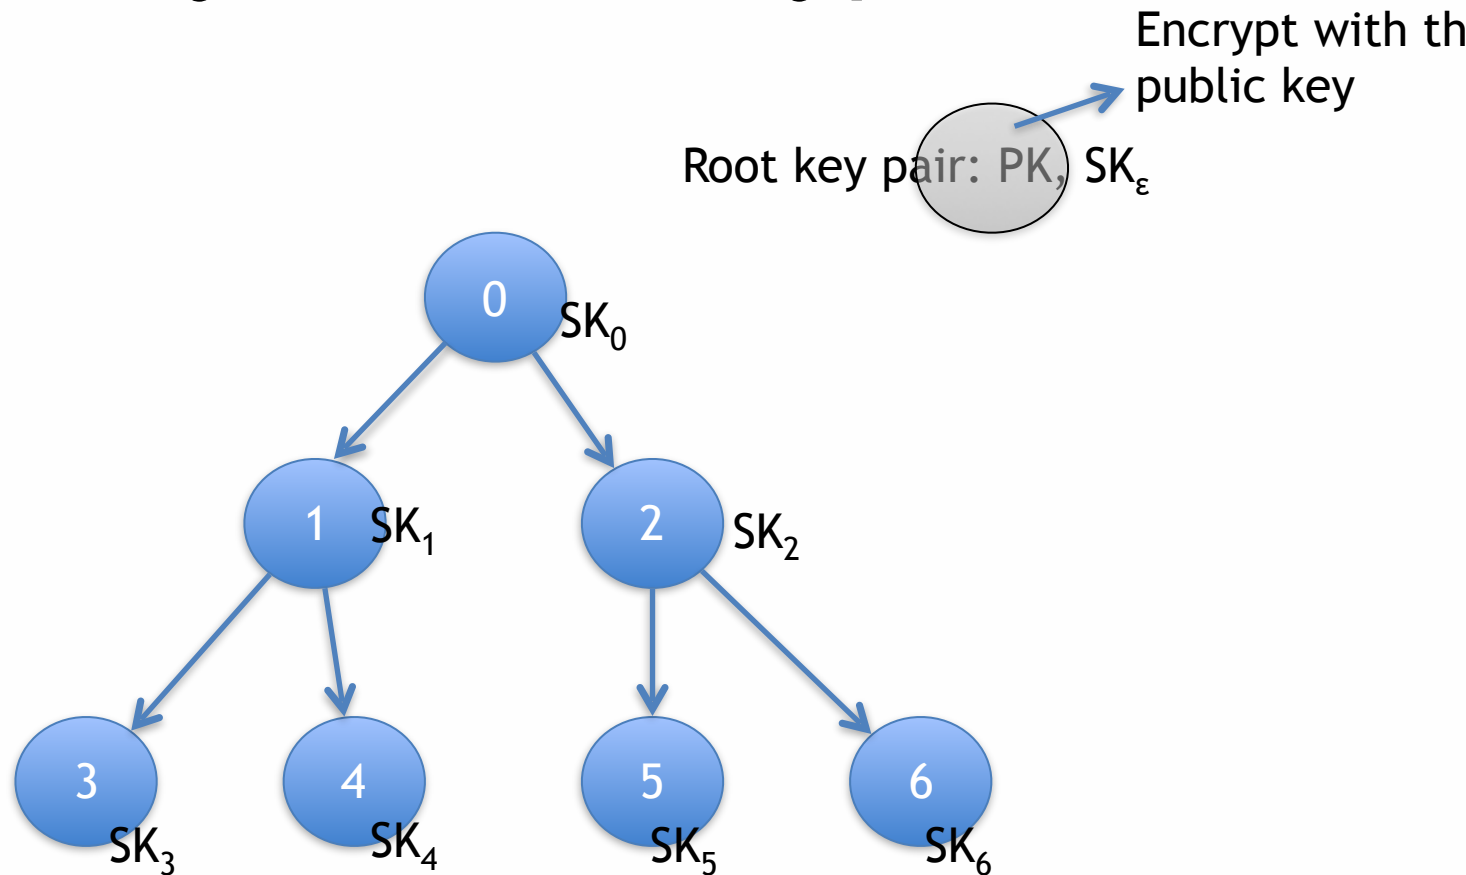
- In plain English, “key-evolving public key encryption”
- Consists of four algorithms:
 - Key generation
 - Key update
 - Encrypt
 - Decrypt

Binary Tree Encryption

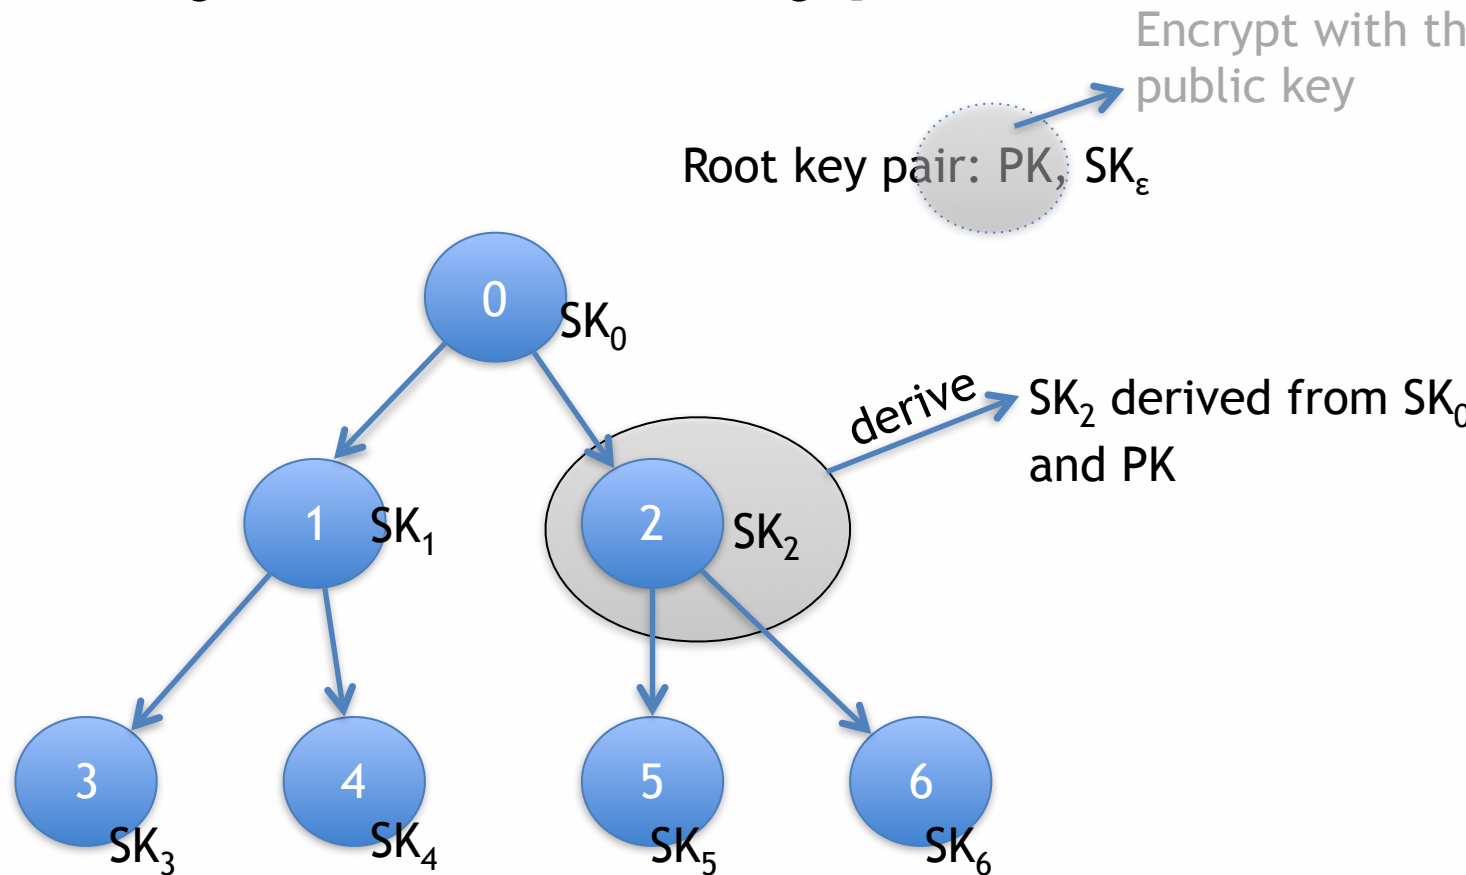
Root key pair: PK, SK_ϵ



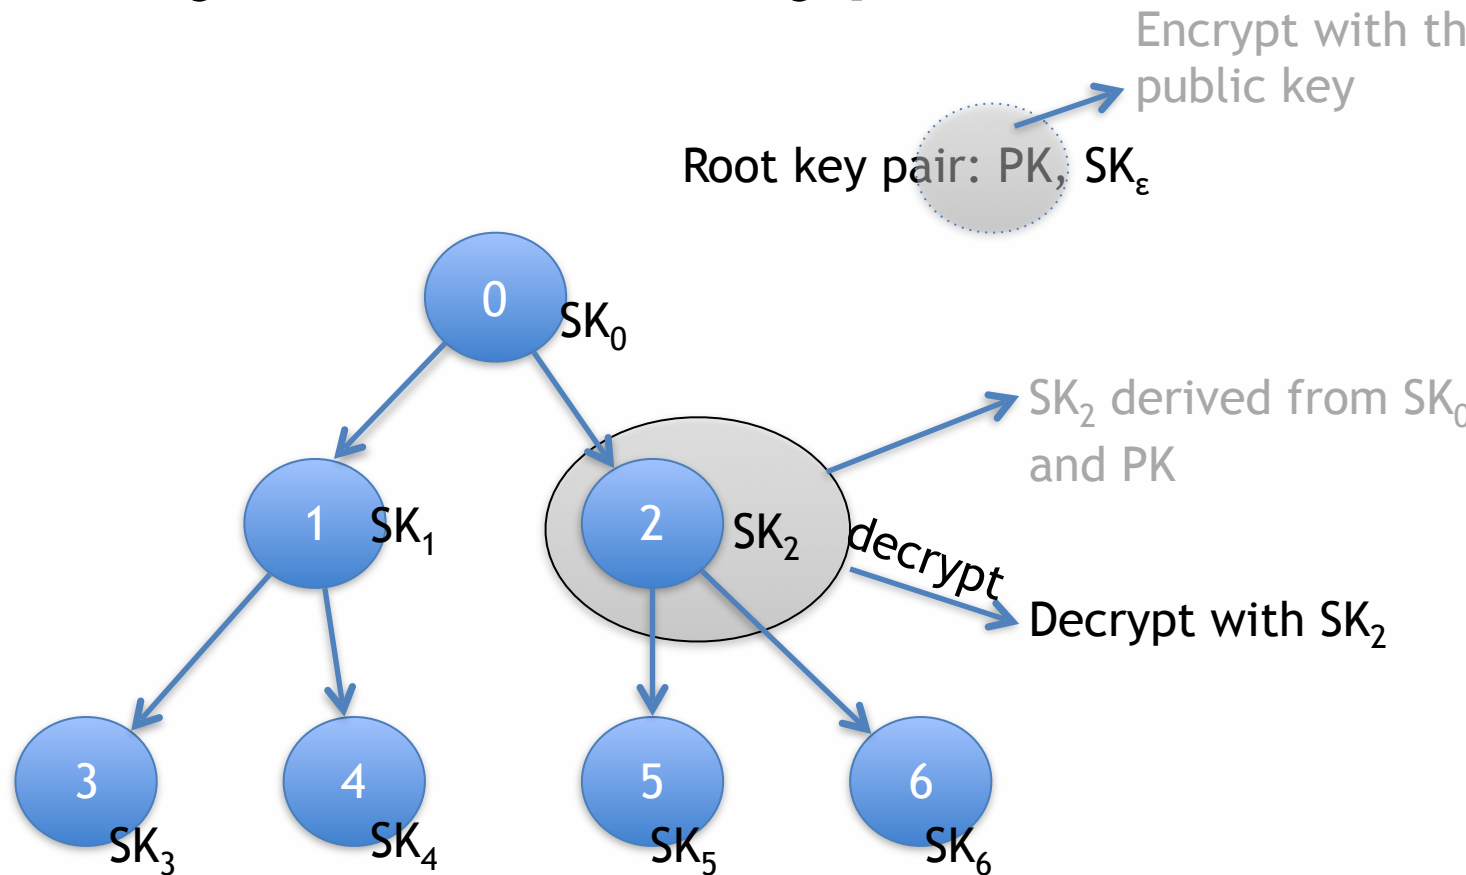
Binary Tree Encryption



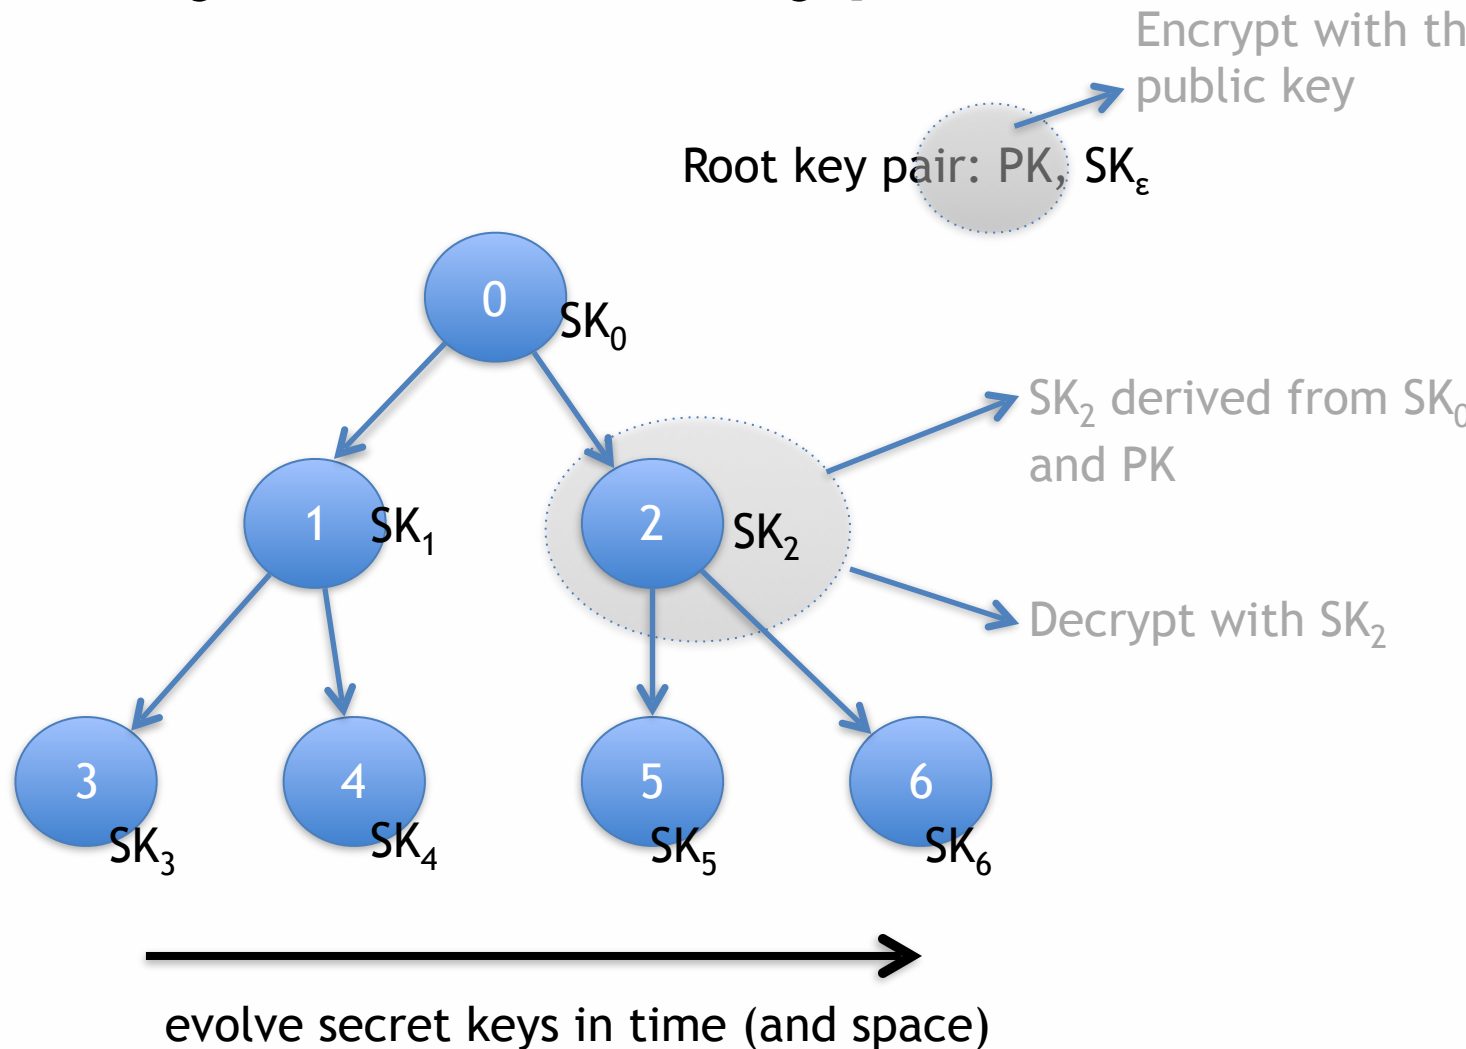
Binary Tree Encryption



Binary Tree Encryption



Binary Tree Encryption



Randomizable Public-Key Encryption

- Many encryption schemes can be re-randomized, e.g., ElGamal, BGN, etc.
- How can we maintain integrity after randomizing ciphertext?
 - *Randomizable signatures*

Randomizable Signatures

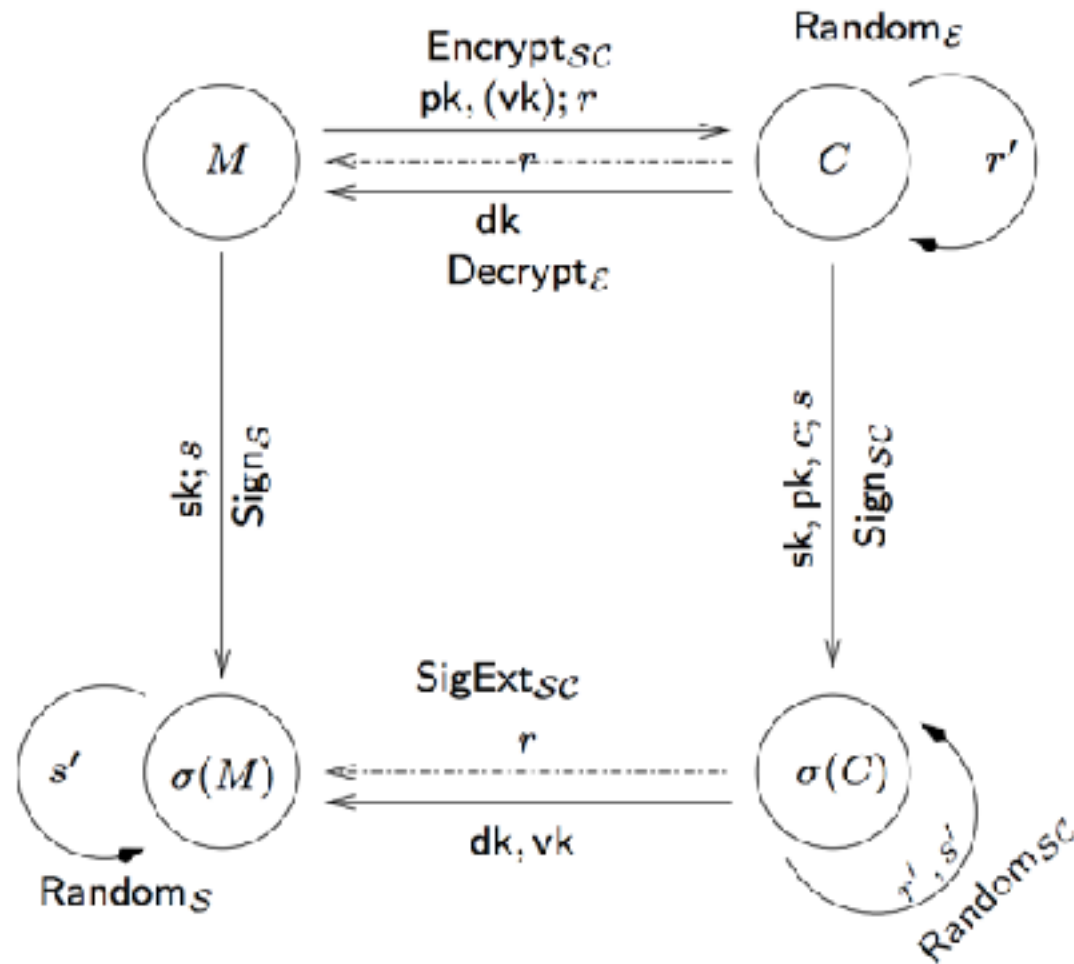
- Given a ciphertext, anyone can re-randomize the ciphertext and adapt the signature the new encryption, i.e.,

$$\mathcal{D}_0 = \{r' \xleftarrow{\$} \mathcal{R}_e; s' \xleftarrow{\$} \mathcal{R}_s : (c' = \text{Encrypt}(\text{pk}, \text{vk}, m; r'), \sigma' = \text{Sign}(\text{sk}, \text{pk}, c'; s'))\}$$

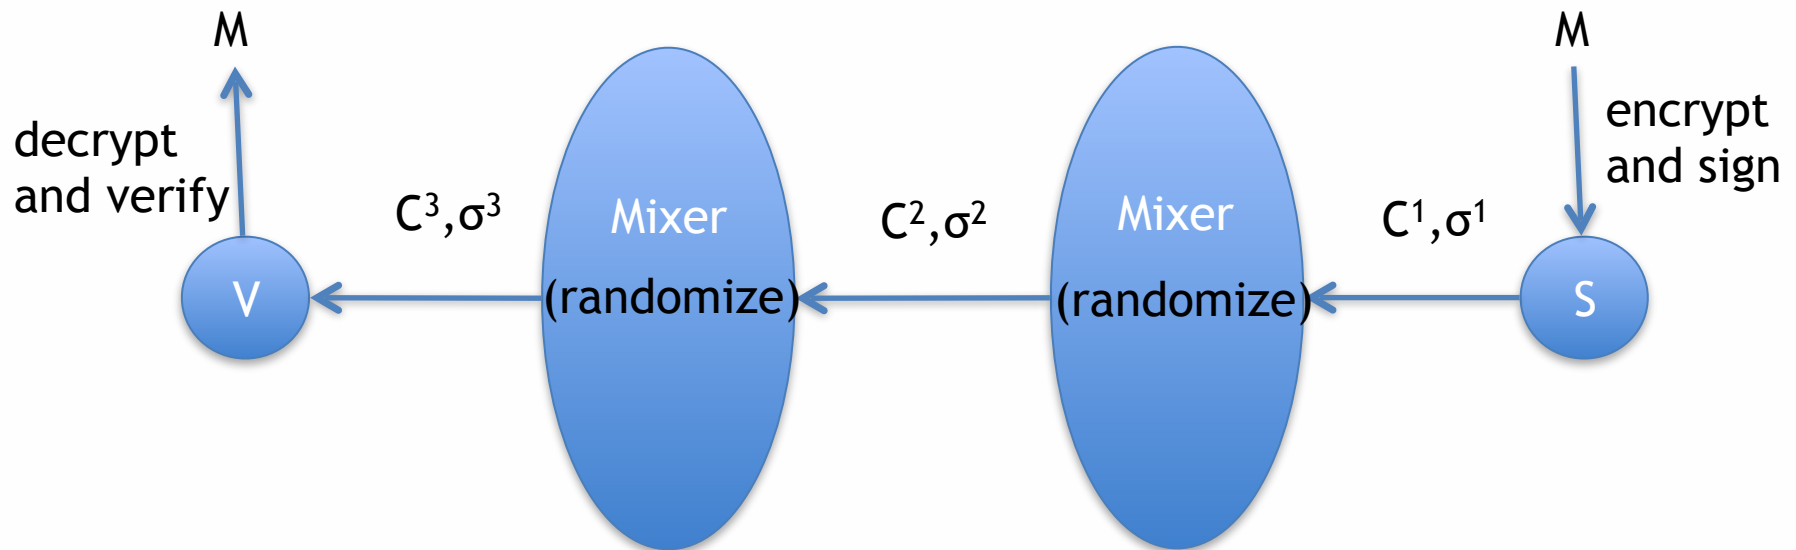
$$\mathcal{D}_1 = \{r' \xleftarrow{\$} \mathcal{R}_e; s' \xleftarrow{\$} \mathcal{R}_s : (c', \sigma') = \text{Random}(\text{vk}, \text{pk}, c, \sigma; r', s')\}$$

are statistically indistinguishable

Extractable Signatures



Mixing in Motion



Privacy: SSE

Searchable Symmetric Encryption

- Goal: search over encrypted data using symmetric-key cryptography
- Many variants:
 - Interactive and non-interactive
 - Response-revealing and response-hiding

Searchable Symmetric Encryption

- Goal: search over encrypted data using symmetric-key cryptography
- Given a database of (encrypted) documents and list of keywords, identify the documents that contain keywords
- Many variants:
 - **Interactive** and non-interactive
 - **Response-revealing** and response-hiding

SSE Overview

- Client possess document list D_1, \dots, D_n
- Client builds an index (DB) that maps keywords w to documents – $DB[w]$
- Protocol mechanics:
 - Setup: return key k and encrypted DB (eDB)
 - Token: on input w and k , return token t
 - Search: on input eDB and t , **the server** returns the identifiers in $DB[w]$
- Complexity: search is at best sub-linear in the number of documents and linear in those containing the keyword (**optimal**)

Other Solutions

Properties	[35, 25]	[35, 25]-light	[40]	[23]	[18]	SSE-1	SSE-2
hides access pattern	yes	yes	no	no	no	no	no
server computation	$O(\log^3 n)$	$O(\sqrt{n})$	$O(n)$	$O(n)$	$O(n)$	$O(1)$	$O(1)$
server storage	$O(n \cdot \log n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$
number of rounds	$\log n$	2	1	1	1	1	1
communication	$O(\log^3 n)$	$O(\sqrt{n})$	$O(1)$	$O(1)$	$O(1)$	$O(1)$	$O(1)$
adaptive adversaries	yes	yes	no	no	no	no	yes

Other Solutions

Properties	[35, 25]	[35, 25]-light	[40]	[23]	[18]	SSE-1	SSE-2
hides access pattern	yes	yes	no	no	no	no	no
server computation	$O(\log^3 n)$	$O(\sqrt{n})$	$O(n)$	$O(n)$	$O(n)$	$O(1)$	$O(1)$
server storage	$O(n \cdot \log n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$	$O(n)$
number of rounds	$\log n$	2	1	1	1	1	1
communication	$O(\log^3 n)$	$O(\sqrt{n})$	$O(1)$	$O(1)$	$O(1)$	$O(1)$	$O(1)$
adaptive adversaries	yes	yes	no	no	no	no	yes

...but the search complexity is always linear in the number of matching documents

Non-Symmetric Encrypted Search

- Two variants: Practical and “not so much”
- Theoretical: Based on ORAM
 - Seek to hide everything except the result size
 - Often multi-round, *stateful*, and have heavy communication costs
- Practical: CryptDB and Arx (CryptDB v2)
 - These are actually deployed

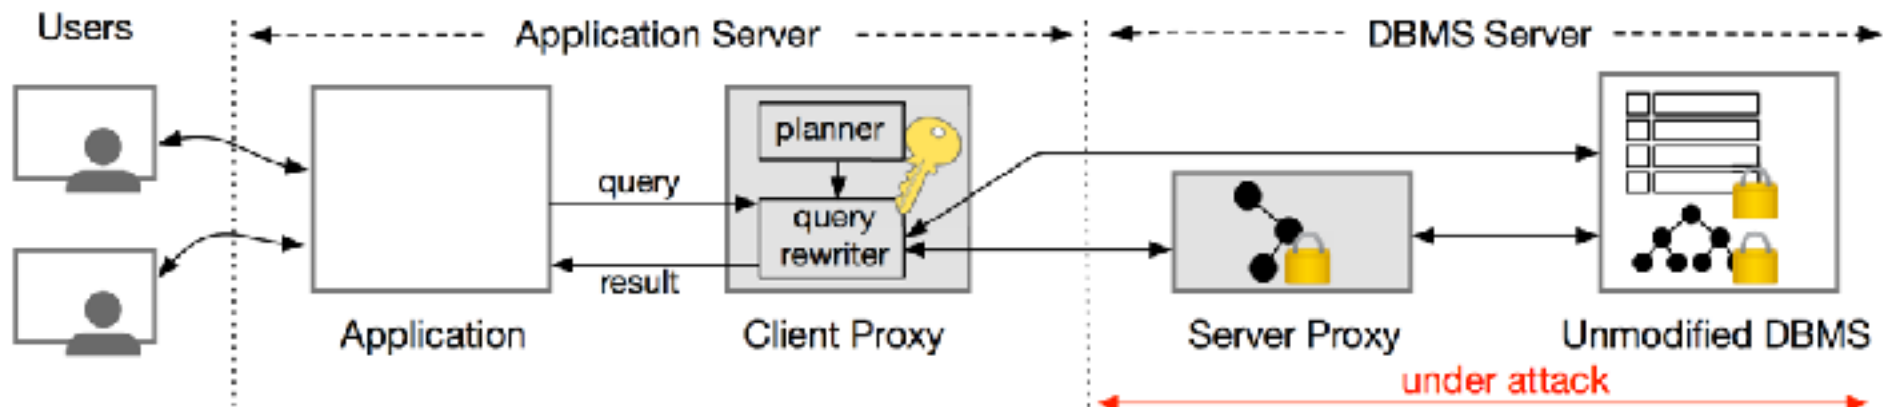
ORAM

Technique	Example Applications	Client-side storage	Blowup
Pointer-based for rooted tree access pattern graph	map/set, priority_queue, stack, queue, oblivious memory allocator	$O(\log N) \cdot \omega(1)$	$O(\log N)$
Locality-based for access pattern graph with doubling dimension \dim	maximum flow, random walk on sparse graphs; shortest-path distance on planar graphs; doubly-linked list, deque	$O(1)^{\dim} \cdot O(\log^2 N) + O(\log N) \cdot \omega(1)$	$O(12^{\dim} \log^{2 - \frac{1}{\dim}} N)$
Path ORAM [45]	All of the above	$O(\log N) \cdot \omega(1)$	$O\left(\frac{\log^2 N + \chi \log N}{\chi}\right)$ for block size $\chi \log N$
ORAM in [29]	All of the above	$O(1)$	$O\left(\frac{\log^2 N}{\log \log N}\right)$

Borrowed from:

<https://eprint.iacr.org/2014/185.pdf>

Arx




Availability

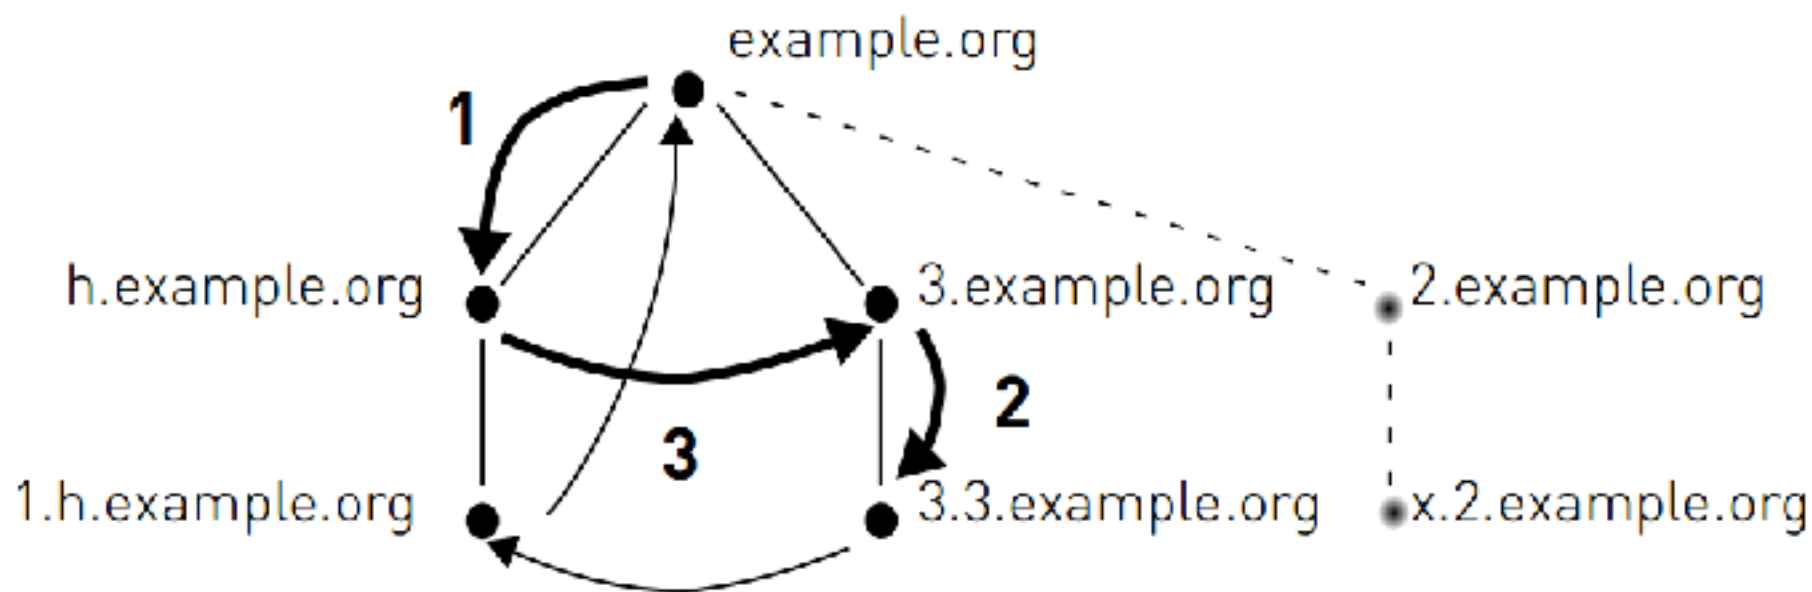
Authenticated DoE

- DNSSEC
 - Prevent forging or modifying DNS records
 - Allow the DNS to prove that a query answer **does not exist**
- Current version allows for zone enumeration

DNSSEC NSEC



DNSSEC NSEC3



NSEC5: Preventing Zone Enumeration

- Main result: public-key operations are necessary to prevent zone enumeration
- Idea: replace the hash in NSEC3 with a keyed hash

Signed Records

Secondary keys: ($PK_S = e$, $SK_S = d$)

For each record x ,

$$S(x) = (h_1(x))^d \bmod N$$

$$F(x) = h_2(S(x))$$

For hash functions h_1 and h_2

Proving Non-Existence

On query q

- Compute $S(q)$ and $F(q)$
- Respond with $S(q)$ and all hashes after $F(q)$

To verify a response

1. Check that $(S(q))^e = h_1(q)$
2. Verify the response with PK_S
3. Check that $F(q) = h_2(S(q))$ is before all hashes

Comments on Confidentiality

- There's a **tremendous number** of papers on confidentiality in ICN
- All of them solve the problem at the application layer
- This problem is **important**, but not yet a core network service