

parc[®]

A Xerox Company

Handling Trust Enforcement

Christopher Wood, PARC

Marc Mosko, PARC

Ersin Uzun, PARC

Agenda

1. Quick refresher on security in CCN
2. Authenticity and trust models
3. Transport-layer enforcement
4. Q&A

Security in CCN

CCN is all about *transferring named data*

All data is named, **secured** (to some degree), and then transferred

Communication security is defined with respect to either (both) of the following items:

Authenticity: ensuring and trusting the legitimacy of content

Confidentiality: controlling who has access to the data

Authenticity and Trust

Why is authenticity important?

- Ensuring consumers get legitimately correct content!
- Prevent network-layer attacks, e.g., ***Content (Cache) Poisoning***

Content Poisoning Attacks

A scenario where a (malicious) entity injects **fake** content into the cache which serves as:

- Invalid responses to consumer interests
- A means to flush legitimate data from the content stores

Authenticity in the Network Layer

Question: How do we protect against attacks at the network layer without knowing what content to trust?

Authenticity in the Network Layer

Question: How to protect against attacks at the network layer without knowing what content to trust?

Answer [1]: Follow the **(Modified) Interest Key Binding rule:**

An Interest message must either reflect the public key of the producer or must be a self-certifying entity.

Implication: include the **hash** of desired content or the **identity of the public verification key**

[1] Ghali, Cesar, Gene Tsudik, and Ersin Uzun. "Network-Layer Trust in Named-Data Networking." ACM SIGCOMM Computer Communication Review 44.5 (2014): 12-19.

Authentication in the Application Layer

Question: How do we determine what content is authentic and trusted?


Authenticity in the Application Layer

Question: How to determine what content is authentic and trusted?

Consider the following **simple** application-layer trust model:

- (A) `lci:/netflix/key` [signed by trusted CA]
- (B) `lci:/netflix/tv/key` [signed by (A)]
- (C) `lci:/netflix/tv/TheOffice` [signed by (B)]
- (D) `lci:/netflix/tv/ParksAndRecreation` [signed by (B)]


Perform signature verification, certificate chain traversal, and check key names against the policy



Fundamental Trust Model Templates

Question: Can we autonomously support any arbitrary trust model?


1. Hierarchical PKI-based model



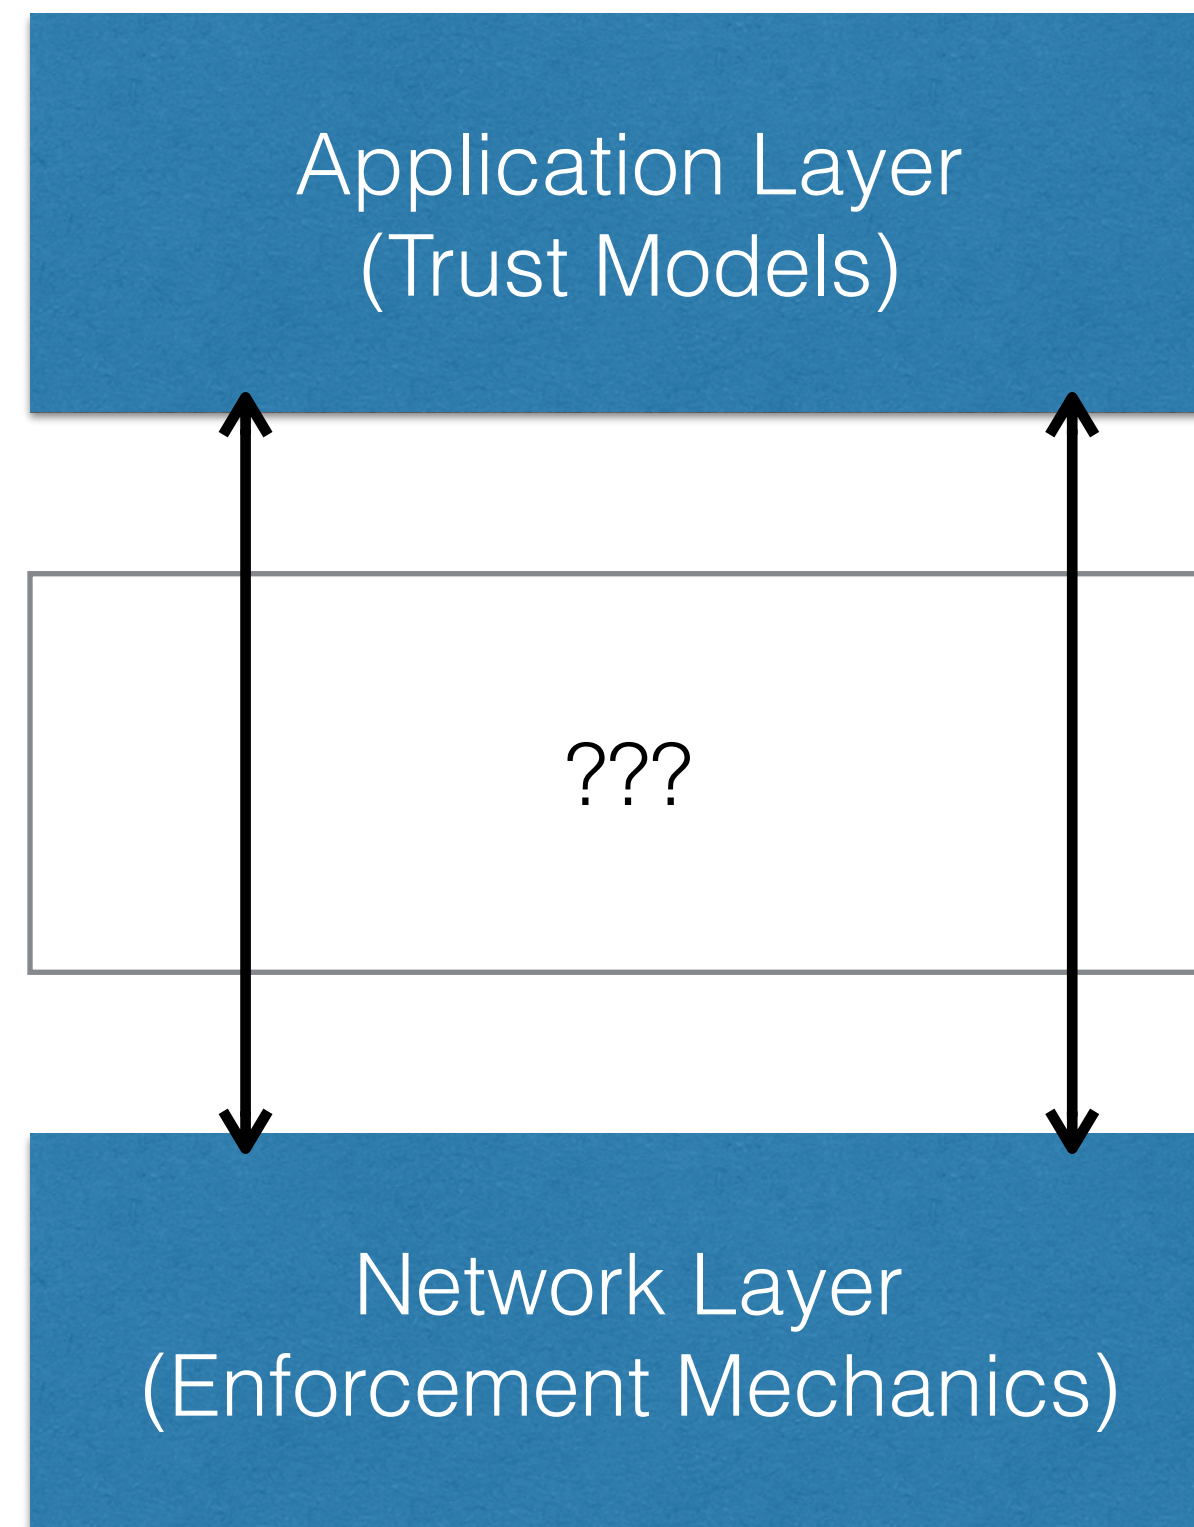
2. Decentralized and distributed
(web of trust)



3. "Flat" and inflexible model

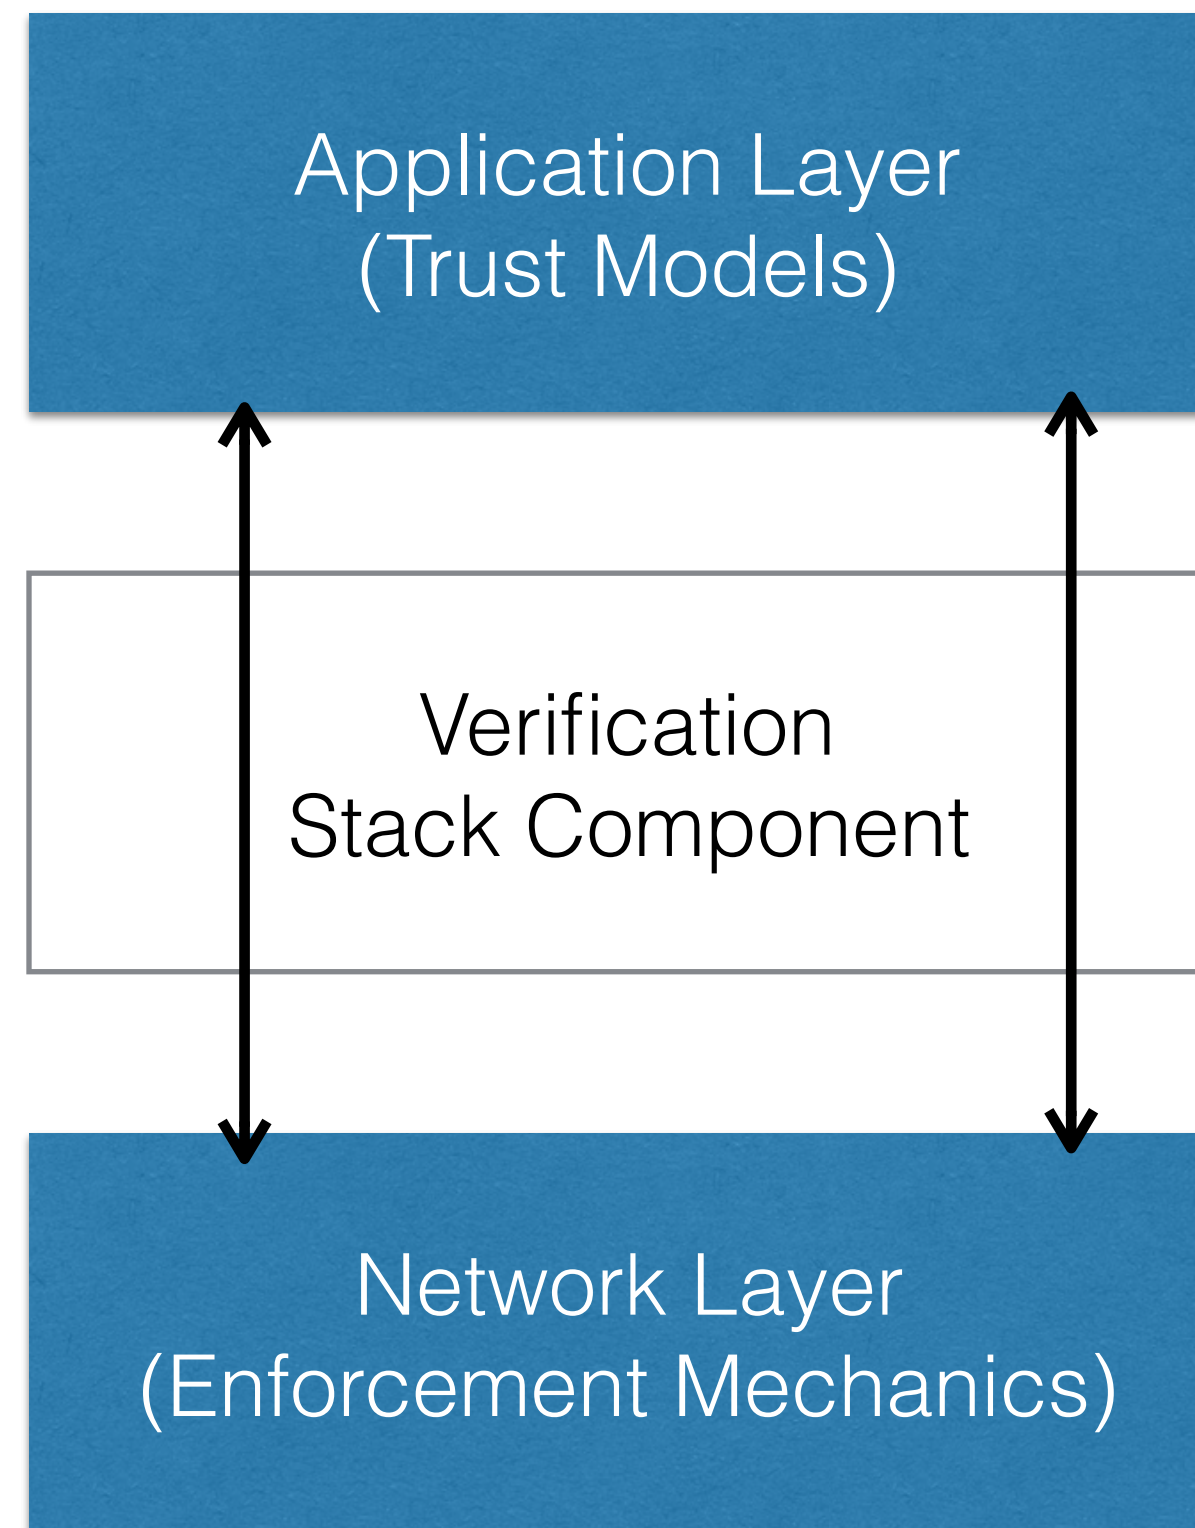


Bridging the Gap

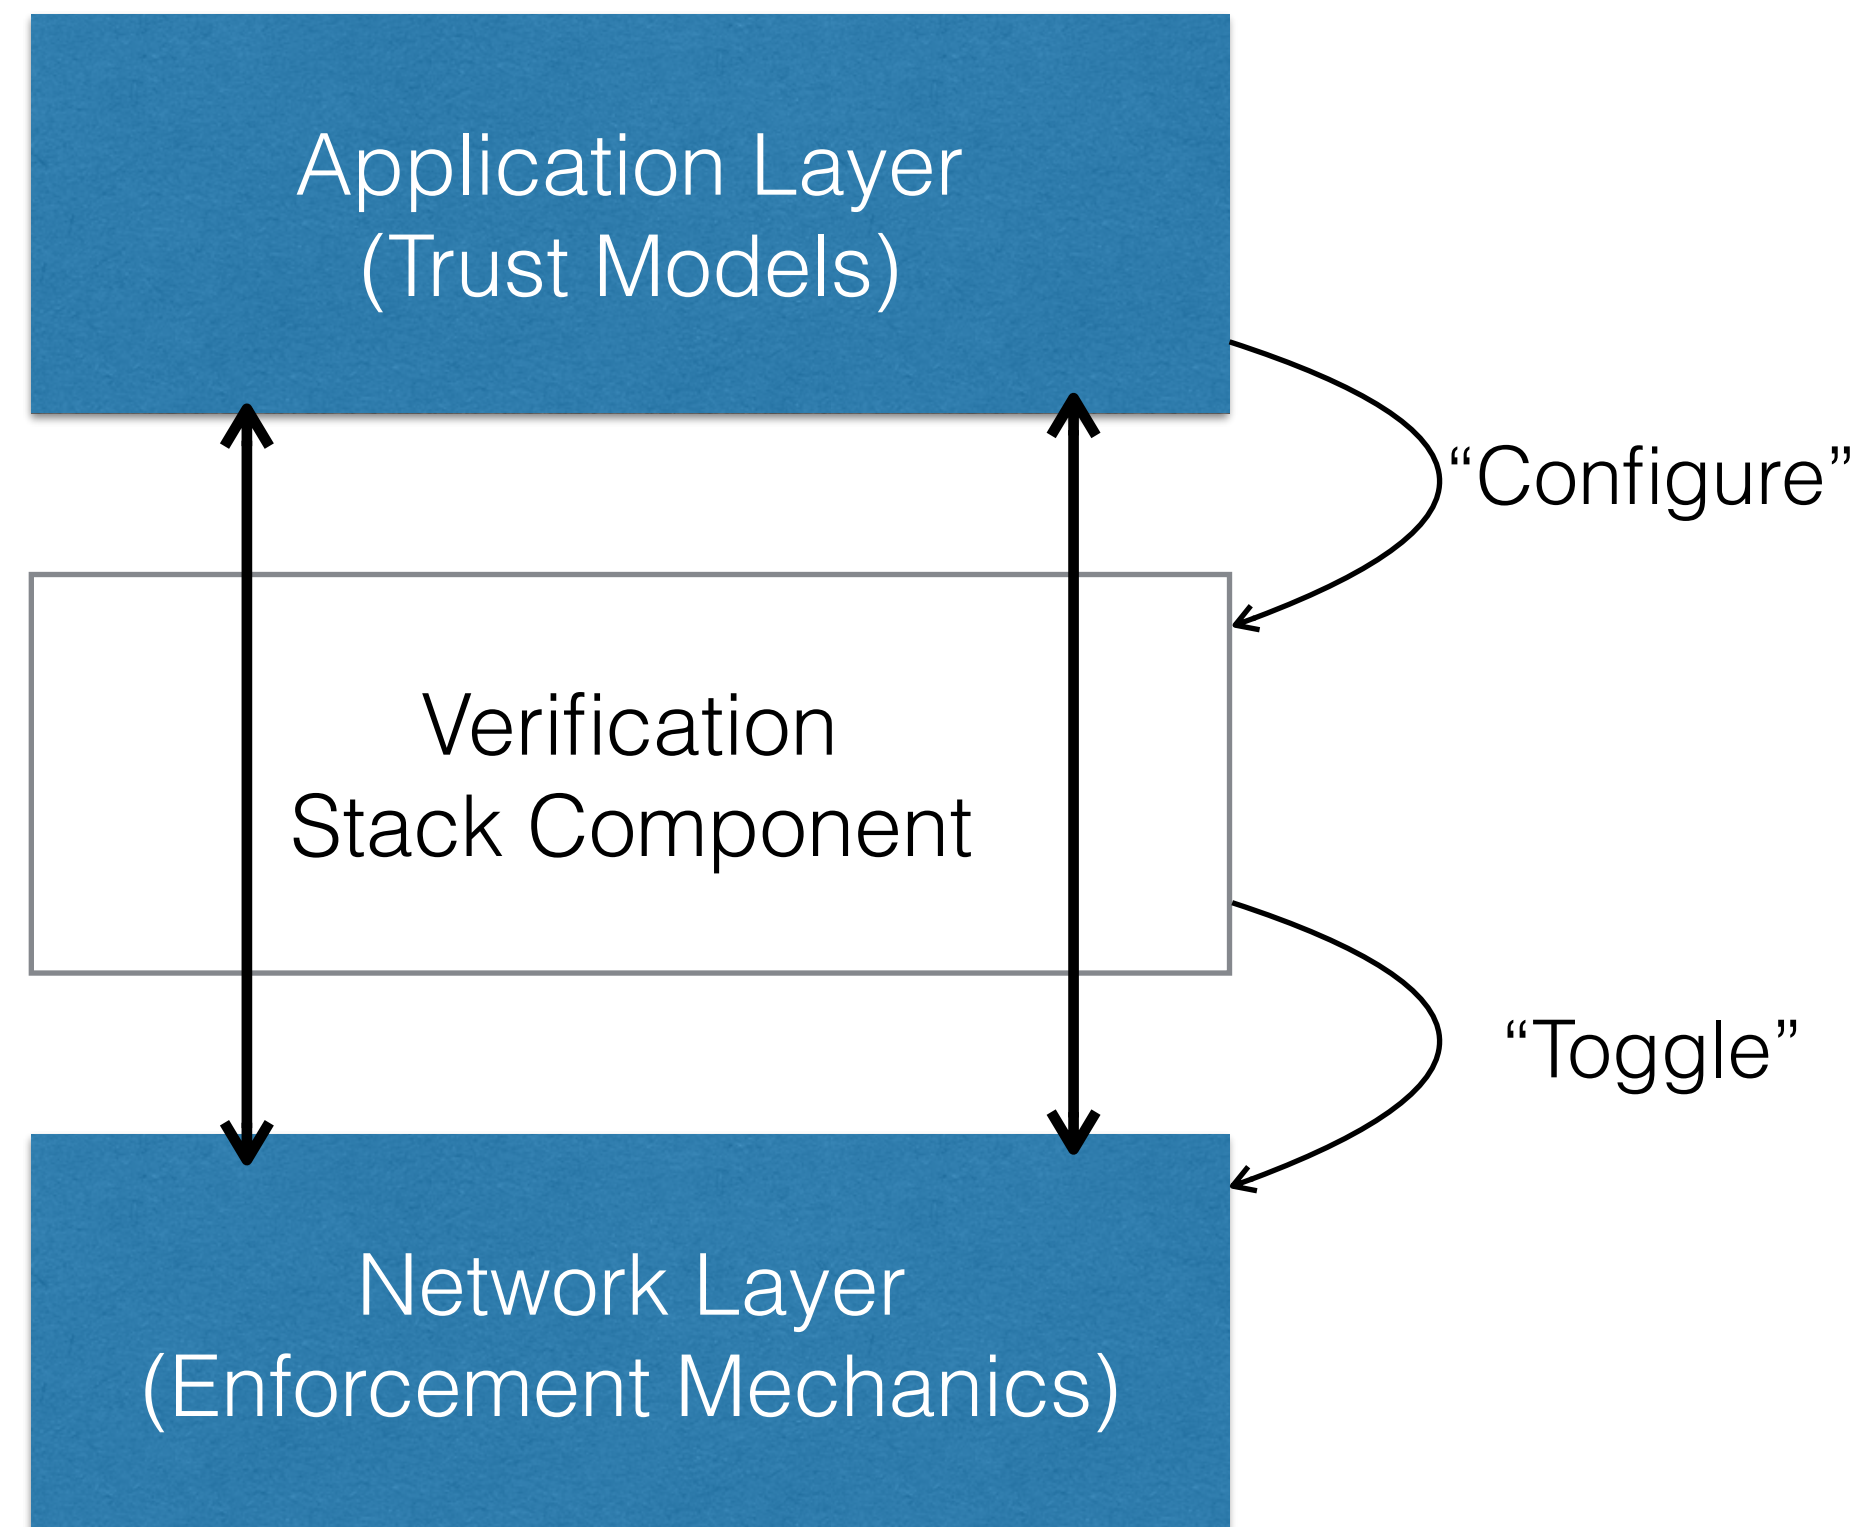


Goal: Enable application-specific trust model instantiation with fundamental trust enforcement machinery at the network layer

Bridging the Gap (Cont'd)



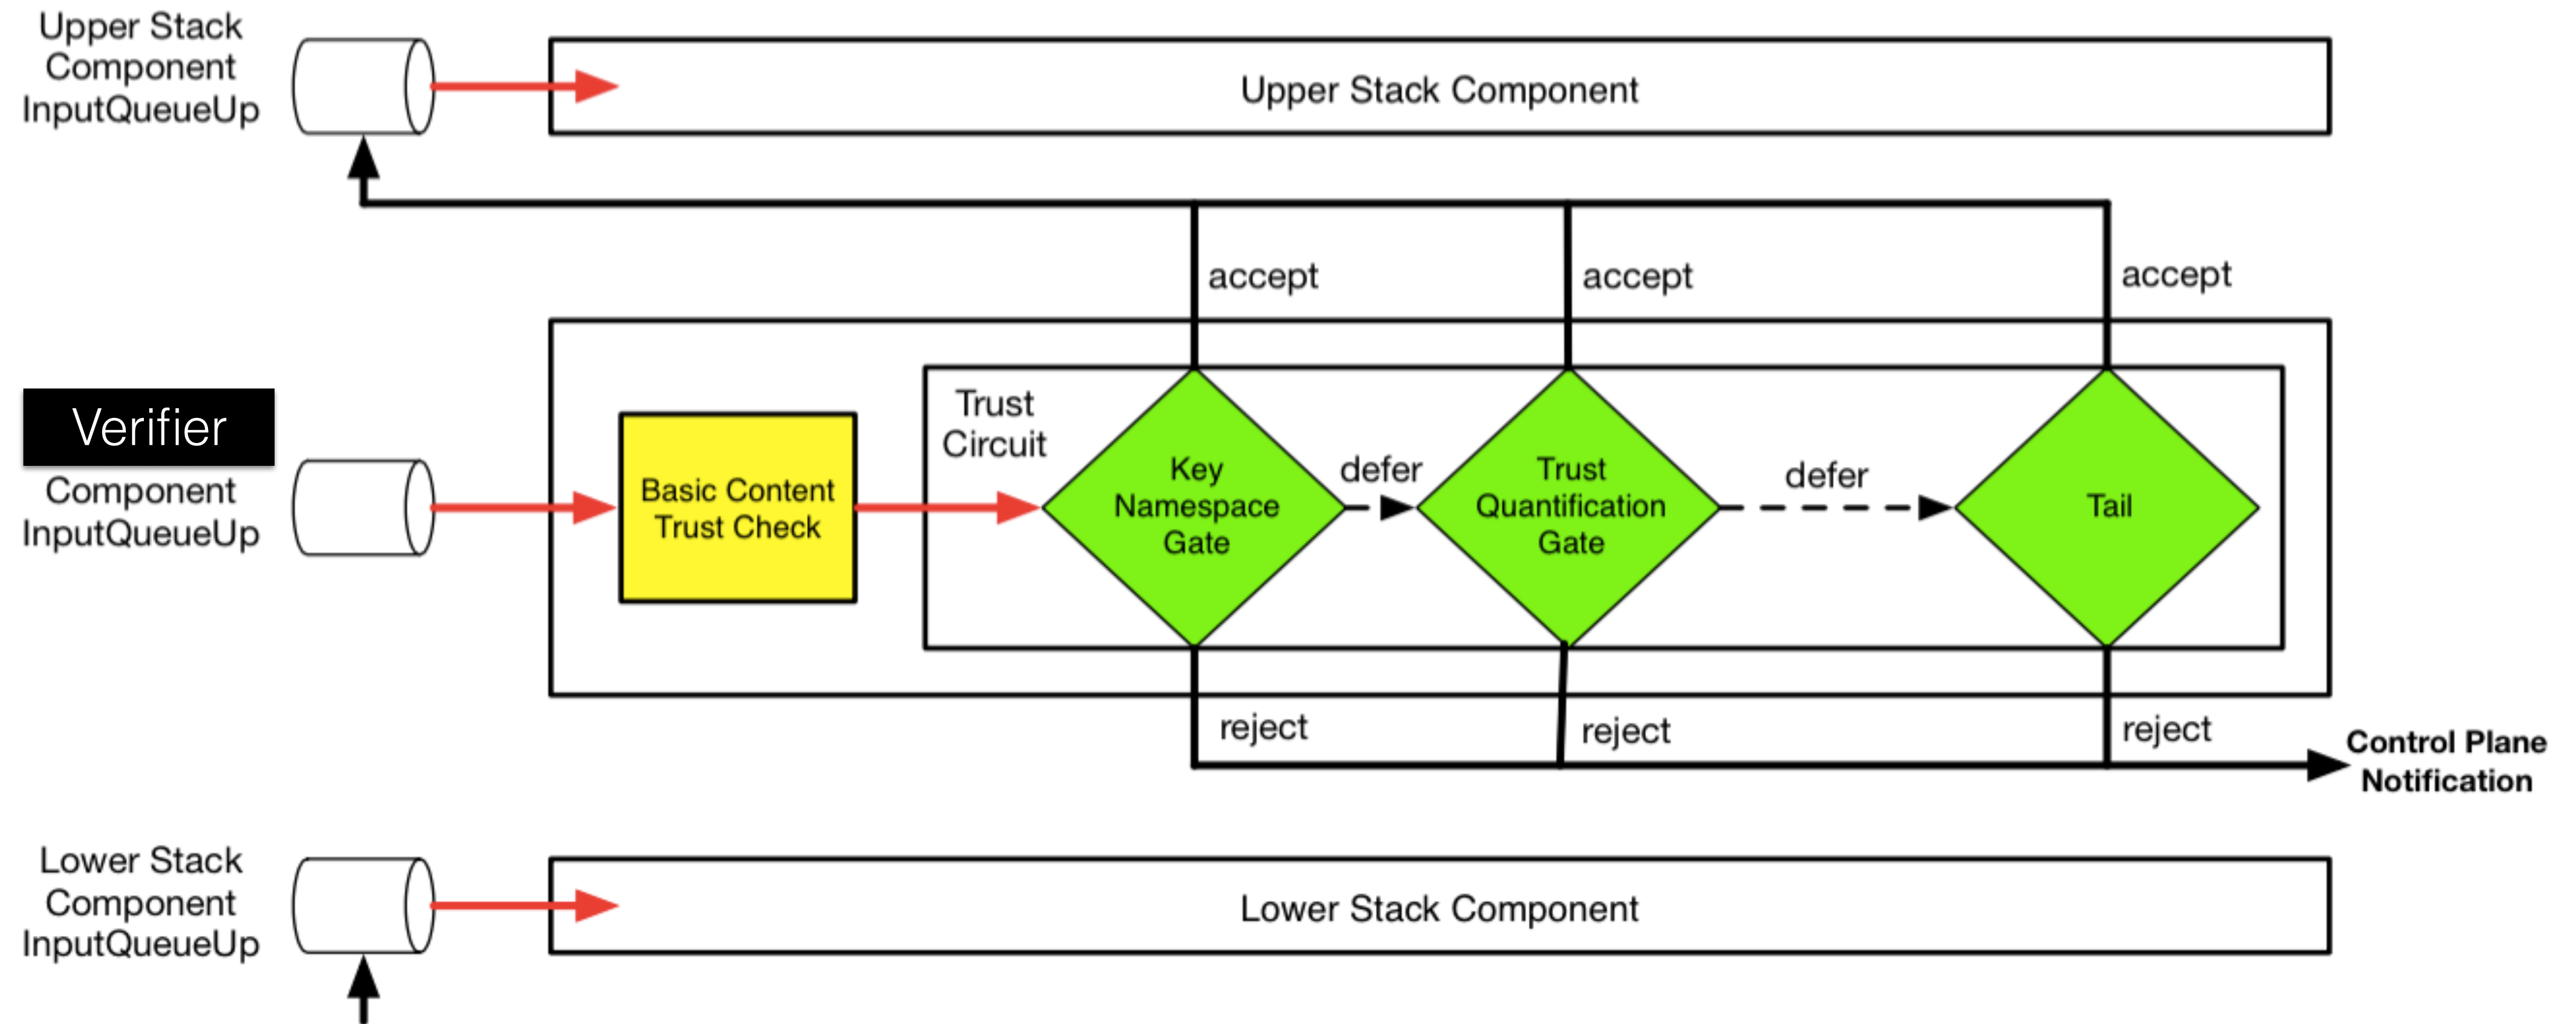
Verification Component



The **Verifier** component is the heart of the trust enforcement machinery responsible for:

1. Verifying ingress content objects
2. Autonomously resolving security information (e.g., certificates)
3. Accepting whitelist and blacklist trust sources
4. ...

Internals: Trust Circuit



Name-Key Binding Trust Models

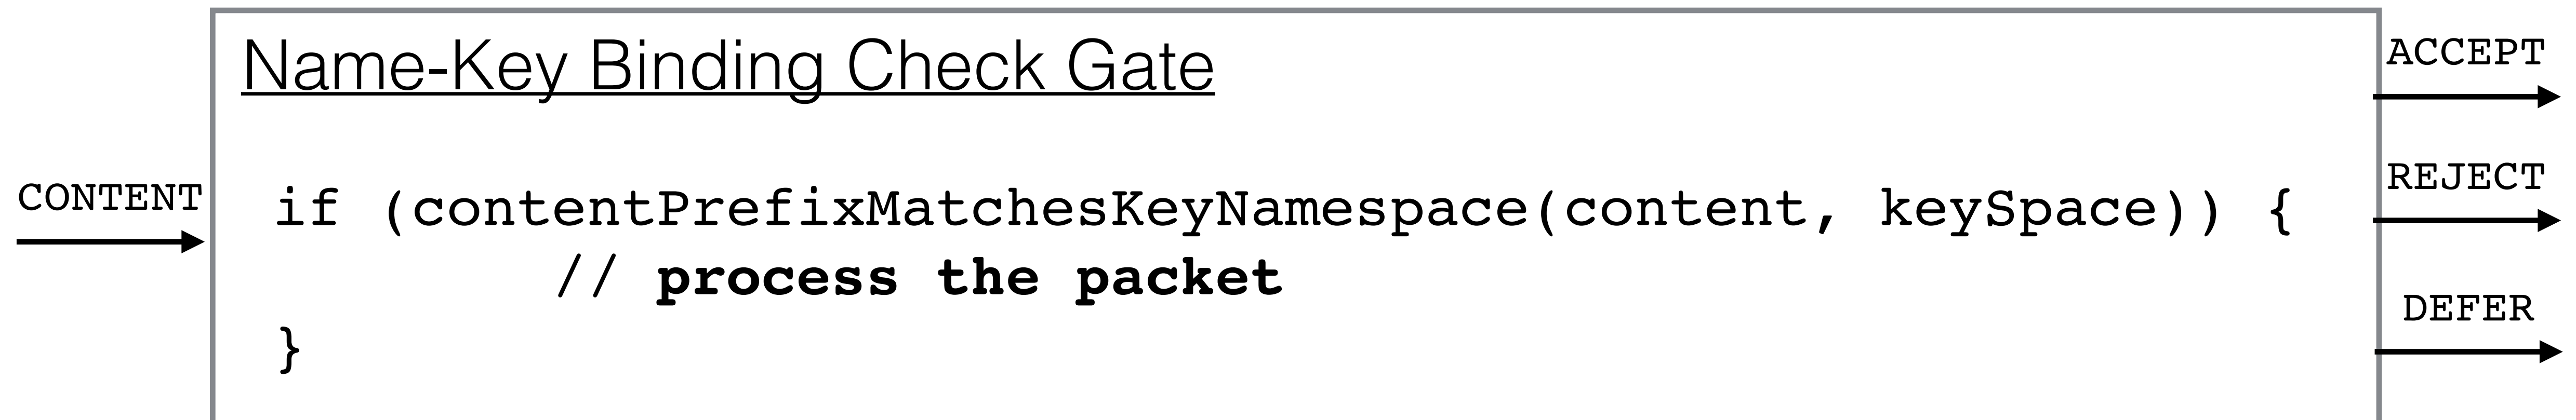
Consider a trust model where **certificates bind keys to name prefixes**

- A content object with name `lci:/parc/csl/file` has a verification key with the name `lci:/xerox/parc/key`
- The content object references the certificate for `lci:/xerox/parc/key`
- The certificate for `lci:/xerox/parc/key` says it can sign content objects with the prefix `lci:/parc/`



Name-Key Binding Trust Model (cont'd)

Enforcement: create a single trust circuit check that checks the certificate-specified prefix against the content object name:



Schematized Binding Trust Model

What if the certificate specified **more than just a name prefix**?

What if the certificate specified **a schema** from which content names are derived?

Content: `lci:/parc/users/cwood/presentations/trust`

Key: `lci:/parc/csl/nds/publickey`

Root key: `lci:/parc/key`

Schema: `[lci:/parc] [/] [users] [/] [*] [/] [presentations | code | ...]`

Producer Implications

Producers sign content according to the expected trust model.

Hierarchical and flat model: use keys issued by trusted CAs or their own trusted parties.

e.g., Mozilla signs Browser plugin content objects with Mozilla's key, which is installed in every user's version of Firefox.

Web of trust: use their own keys (since they are the "introducer")

Consumer Implications

Consumers verify content according to their specified trust models.

The Verifier configuration determines what trust model is to be enforced:

Flat: instantiate the Verifier component with trusted (whitelisted) roots and disable certificate chain traversal


Hierarchical: instantiate the Verifier component with trusted (whitelisted) roots and enable certificate chain traversal

...

Questions?...

parc[®]

A Xerox Company



Thank you